


# Intro to MIC performance

Vasily Volkov

March 6, 2012







# Agenda

- Three programming models for MIC
- Hitting best throughput numbers
- OpenMP thread affinity settings
- Reproducing high performance in MKL
- Performance of cache coherence mechanism
- Beating Intel OpenMP barrier implementation

# **Part I**

## Basics of MIC programming

# Starting up

I. SSH to MIC systems at:

- **`knight.millennium.berkeley.edu`**
  - 32 cores @ 1.2 GHz
  - Used throughout this presentation
- **`ferry.millennium.berkeley.edu`**
  - 30 active cores @ 1.05 GHz

II. Set up compiler environment:

**`source /opt/intel/composerxe_mic/bin/compilervars.sh intel64`**

# Three ways to program MIC

## Native mode

- SSH to MIC and use it as any multiprocessor machine

## (Heterogeneous) explicit copy model

- Mark up the code that runs on MIC using #pragma
- Resembles OpenMP and CUDA
- Specify which arrays to transfer to MIC and back

## (Heterogeneous) implicit copy model

- Mark up calls to MIC using keywords
- All transfers to MIC are handled automatically

# Computing $\pi$ in MIC native mode

```
#include <stdio.h>
main()
{
    double sum = 0, h = 1./1024/1024;
    for( double x = h/2; x < 1; x += h )
        sum += 4/(1+x*x);
    printf( "pi = %.15f\n", sum*h );//15 accurate digits
}
```

At host command line:

**icc -mmic pi.cpp**

**sudo scp -i /opt/intel/mic/id\_rsa a.out root@192.168.1.100:**

**sudo ssh -i /opt/intel/mic/id\_rsa root@192.168.1.100**

At MIC command line: **./a.out**

Ask support@millennium for sudo access

# Computing $\pi$ in MIC offload mode

```
#include <stdio.h>
main()
{
    double sum = 0, h = 1./1024/1024;
    #pragma offload target(mic)
    {
        for( double x = h/2; x < 1; x += h )
            sum += 4/(1+x*x);
    }
    printf( "pi = %.15f\n", sum*h );
}
```

Host command line:

**icc -offload-build pi.cpp**

**./a.out**

Note: if MIC is unavailable, everything runs on host



# Computing $\pi$ in MIC Cilk+ mode

```
#include <stdio.h>

_Cilk_shared double compute_pi( double h )
{
    double sum = 0;
    for( double x = h/2; x < 1; x += h ) sum += 4/(1+x*x);
    return sum*h;
}

main()
{
    double pi = _Cilk_offload compute_pi( 1./(1<<20) );
    printf( "pi = %.15f\n", pi );
}
```

At host command line:

**icc -offload-build pi.cpp**

**./a.out**

# Explicit copy model

```
#include <stdio.h>
__declspec(target(mic)) double compute( double *x, int n )
{
    double sum = 0;
    for( int i = 0; i < n; i++ ) sum += 4/(1+x[i]*x[i]);
    return sum/n;
}
main()
{
    int n = 1<<20;
    double *x = (double *)malloc( n*sizeof(double) ), pi;
    for( int i = 0; i < n; i++ ) x[i] = (i+0.5)/n;
    #pragma offload target(mic) in(x:length(n))
        pi = compute( x, n );
    printf( "pi = %.15f\n", pi );
}
```

To transfer an array to MIC use **in(...)**

Other options: out(...), inout(...), nocopy(...)

# Implicit copy model

```
#include <stdio.h>
_Cilk_shared double compute( double _Cilk_shared *x, int n )
{
    double sum = 0;
    for( int i = 0; i < n; i++ ) sum += 4/(1+x[i]*x[i]);
    return sum/n;
}
main()
{
    int n = 1<<20;
    double _Cilk_shared *x = (double _Cilk_shared*)
        _Offload_shared_malloc( n*sizeof(double) );
    for( int i = 0; i < n; i++ ) x[i] = (i+0.5)/n;
    double pi = _Cilk_offload compute( x, n );
    printf( "pi = %.15f\n", pi );
}
```

Automatic transfers if using **\_Offload\_shared\_malloc** and **\_Cilk\_offload**

# Recommended Reading

<https://mic-dev.intel.com>

- Need Intel contact to get access
- Of interest: architecture specs, forum

[/opt/intel/composerxe\\_mic/Documentation/en\\_US/mic/mic\\_compiler\\_guide.pdf](#)

- Programming models, language extensions

[/opt/intel/composerxe\\_mic/Documentation/en\\_US/Release\\_Notes\\_mic.pdf](#)

- Section 4.4.21: Using large pages in DMA buffers
- Section 4.4.22: Using large pages in malloc

# A few more system tools

System info on MIC: **sudo micinfo**

Software stack is still not very stable...

- To reboot MIC only: **sudo micstart -r**
- To reboot entire system: **sudo reboot**



## **Part II**

### Measuring offload overhead

# Offload overhead in pragma model

```
#include <stdio.h>
#include <mkl.h>
main( )
{
    for( int i = 0; i < 100; )
    {
        double t = dsecnd(); //accurate timer from MKL
        #pragma offload target(mic)
        {
            i++;
        }
        t = dsecnd() - t;
        printf( "%g ms\n", t*1e3 );
    }
}
```

Compile this **and all following codes** using:  
**icc -offload-build -O3 -mkl bench.cpp**

# Offload overhead in Cilk+ model


```
#include <stdio.h>
#include <mkl.h>

_Cilk_shared int foo( int i ) { return i+1; }

main( )
{
    for( int i = 0; i < 100; )
    {
        double t = dsecnd();
        i = _Cilk_offload foo( i );
        t = dsecnd() - t;
        printf( "%g ms\n", t*1e3 );
    }
}
```

Compile this **and all following codes** using:  
**icc -offload-build -O3 -mkl bench.cpp**

# Offload overhead



# Large overheads in Cilk+ are going to stay

Kevin Davis (Intel):

Development does not consider the `_Cilk_shared` behavior [40+ ms overhead] to be a defect...

<https://mic-dev.intel.com/node/1508#comment-3161>



## **Part III**

### Arithmetic peak

# Getting the arithmetic peak (1/2)

```
#ifdef __MIC__ //can't use intrinsics if not on MIC
#include <micvec.h>
__declspec(target(mic)) void foo( float &r, int niterations )
{
    F32vec16 a( 100.f ), c( 1.f ), x( 0.9f );
    for( int i = 0; i < niterations; i += 8 )
    {
        a = a*x + c; a = a*x + c; a = a*x + c; a = a*x + c;
        a = a*x + c; a = a*x + c; a = a*x + c; a = a*x + c;
    }
    r = reduce_add( a );
}
#else
__declspec(target(mic)) void foo( float &r, int niterations ){}
#endif
```

# Getting the arithmetic peak (2/2)

```
#include <stdio.h>
#include <mk1.h>
main( )
{
    for( int i = 0; i < 20; i++ )
    {
        int niterations = (100+i)<<20;
        float t, r;
        #pragma offload target(mic)
        {
            double t0 = dsecnd();
            foo( r, niterations );
            t = dsecnd() - t0;
        }
        printf( "%g Gflop/s, %g\n", 32e-9*niterations/t, r );
    }
}
```

# Arithmetic throughput results

Result: 9.5 Gflop/s = **25% of 1-core peak**

Why: MAD back-to-back latency is 4 cycles

- **Need 4 MADs in the flight to get the peak**

Issues from same thread only every other cycle

- **Need at least 2 threads to get the peak**

# NOPs?!

Disassemble: **icc -offload-build -O3 -mkl -S bench.cpp**

```
..B2.3:                                # Preds ..B2.1 ..B2.3 Latency 29
    vmadd213ps %v1, %v0, %v2            #8.19 c1
    addl      $8, %eax                  #6.38 c1
    cmpl      %esi, %eax                #6.25 c3
    vmadd213ps %v1, %v0, %v2            #8.32 c5
    nop                                #8.45 c7
    vmadd213ps %v1, %v0, %v2            #8.45 c9
    nop                                #8.58 c11
    vmadd213ps %v1, %v0, %v2            #8.58 c13
    nop                                #9.19 c15
    vmadd213ps %v1, %v0, %v2            #9.19 c17
    nop                                #9.32 c19
    vmadd213ps %v1, %v0, %v2            #9.32 c21
    nop                                #9.45 c23
    vmadd213ps %v1, %v0, %v2            #9.45 c25
    nop                                #9.58 c27
    vmadd213ps %v1, %v0, %v2            #9.58 c29
    jl        ..B2.3                    # Prob 82%    #6.25 c29
```

**Can't get >50% of peak with this code (NOPs don't pair)**



# Why NOPs?

Most hazards on MIC stall entire core, not a thread:

- Instruction cache miss
- L2 cache miss (but not prefetch)
- Full prefetch queue
- Vector register dependency
- Cache bank conflicts between U and V pipe
- Misaligned access
- Forwarding from store queue to load unit

**Use NOP/DELAY to pause a thread to avoid core stall**

- Compiler tends to add NOPs automatically
- It can be disabled: `-offload-copts="-mCG_lrb_num_threads=4"`
- (unofficial option, may be dropped in the future)

# Two MADs in flight: 49% peak (1 core)

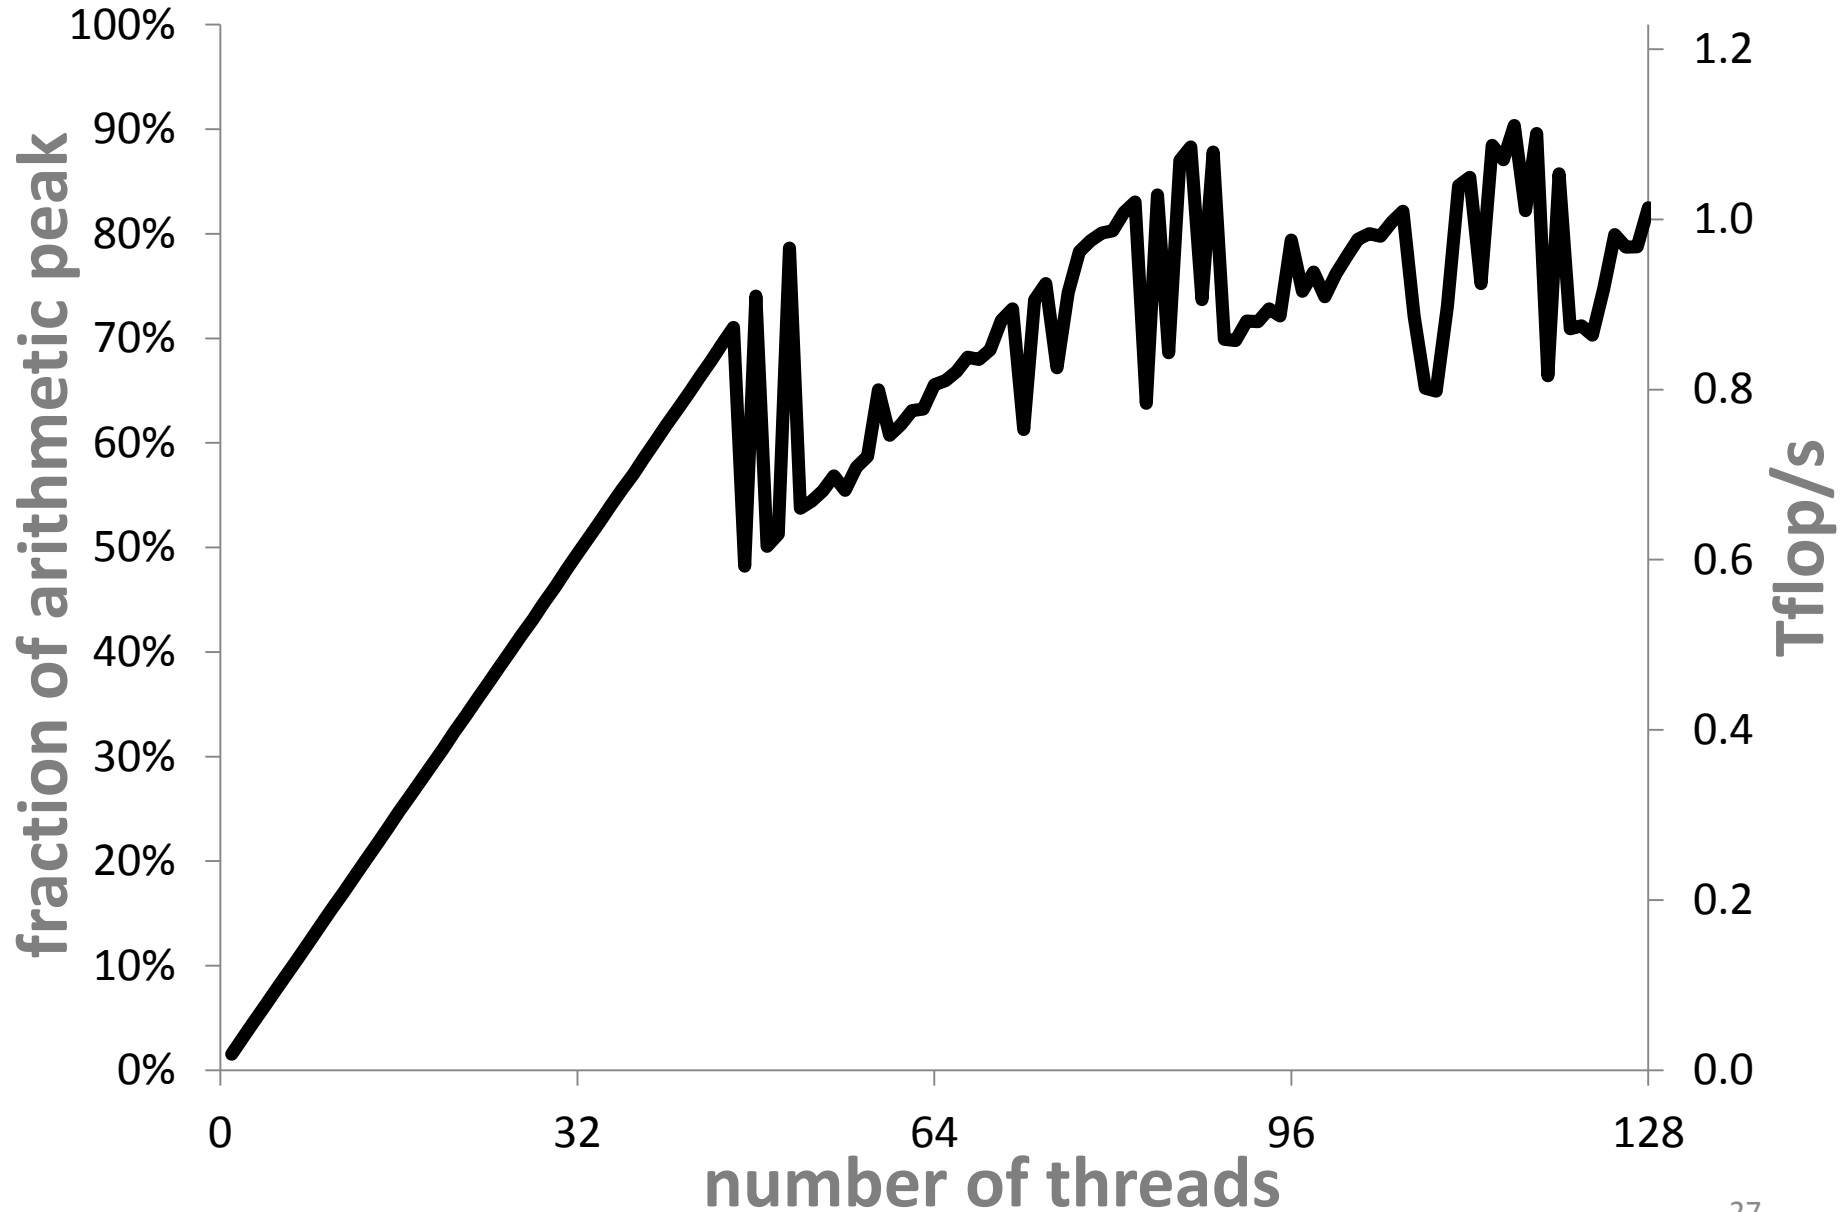
```
#ifdef __MIC__ //can't use intrinsics if not on MIC
#include <micvec.h>
__declspec(target(mic)) void foo( float &r, int iterations )
{
    F32vec16 a( 100.f ), b( 200.f ), c( 1.f ), x( 0.9f );
    for( int i = 0; i < iterations; i += 8 )
    {
        a = a*x + c; b = b*x + c; a = a*x + c; b = b*x + c;
        a = a*x + c; b = b*x + c; a = a*x + c; b = b*x + c;
    }
    r = reduce_add(a+b);
}
#else
__declspec(target(mic)) void foo(float &r, int iterations){}
#endif
```

No NOP problem here as no back-to-back dependencies


# Run many threads using OpenMP

```
#include <stdio.h>
#include <omp.h>
#include <mk1.h>
main()
{
    float Gflops, r[1000];
    for( int i = 0; i < 128; i++ )
    {
        int niterations = 100<<20, nthreads = ((i*79)%128)+1;
        #pragma offload target(mic)
        for( int k = 0; k < 3; k++ )//first two runs are warm up
        {
            double t = dsecnd();
            #pragma omp parallel num_threads(nthreads)
            foo( r[omp_get_thread_num()], niterations );
            Gflops = 32e-9*niterations*nthreads/(dsecnd()-t);
        }
        printf( "%d threads, %g Gflop/s, %g\n", nthreads, Gflops, r[0]);
    }
}
```

# How throughput grows with #threads

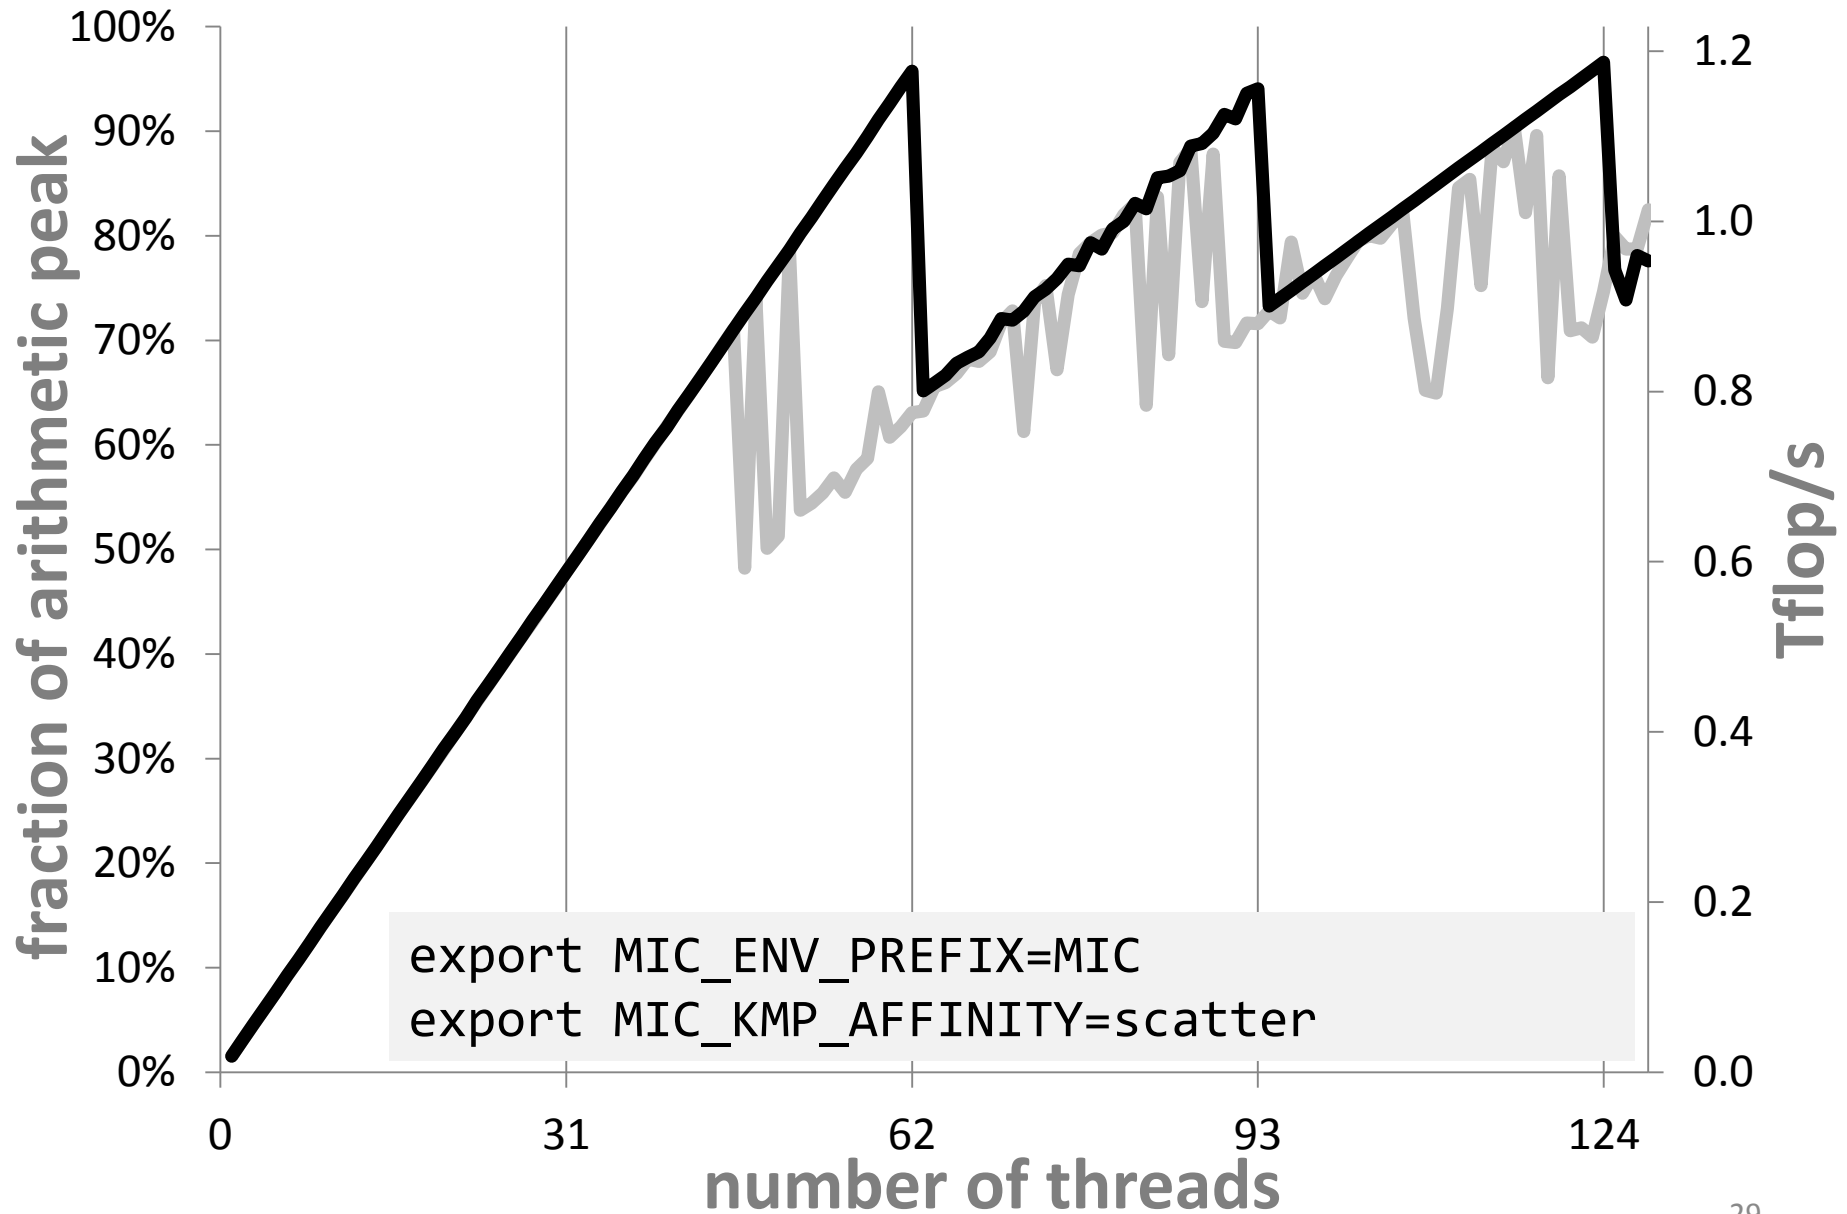


# Enforce round-robin thread assignment






# One core is reserved in offload mode




# Enable the reserved core



# Common KMP\_AFFINITY settings

128 logical processors



`norespect`: enables core #31 (invisible otherwise)  
`compact`: assign in the order 1,2,3,4,...,124,0,125,...  
`scatter`: round-robin order – 1,5,9,...,121,2,6,...  
`explicit,proclist=[...]`: any user-defined order  
`granularity=fine`: disable thread migration within core

# Figuring out thread affinity

**export MIC\_KMP\_AFFINITY=...,verbose**

- Reports mapping of OpenMP threads to OS threads to hardware thread contexts

**sched\_getcpu() / #include <sched.h>**

- Returns current OS thread

```
#define cpuid_getcpu() ({unsigned int eax,ebx;__asm__\  
("cuid": "=a"(eax), "=b"(ebx): "0"(1): "cx", "dx");ebx>>24;})
```

- $\text{cpuid\_getcpu()}/4$  = physical core
- $\text{cpuid\_getcpu()} \% 4$  = local thread context

## **Part IV**

### Memory peak

# Prefetching on MIC

Chip pin bandwidth: 256 pins@1.5GHz = 96 GB/s

- Need 30 KB in flight to hide 300 cycle latency
- This is over 200 B/thread if using 128 threads
- Cache line size is 64 B

So, multithreading alone is not enough

- But entire core stalls on an L2 miss anyway

**Must use software prefetching** (no HW prefetch)

- Inserted by compiler or manually
- `#pragma noprefetch` disables compiler prefetching
- Page fault if prefetching invalid pages
  - **Allocate extra trailing space**

# Manual prefetching example (1/2)

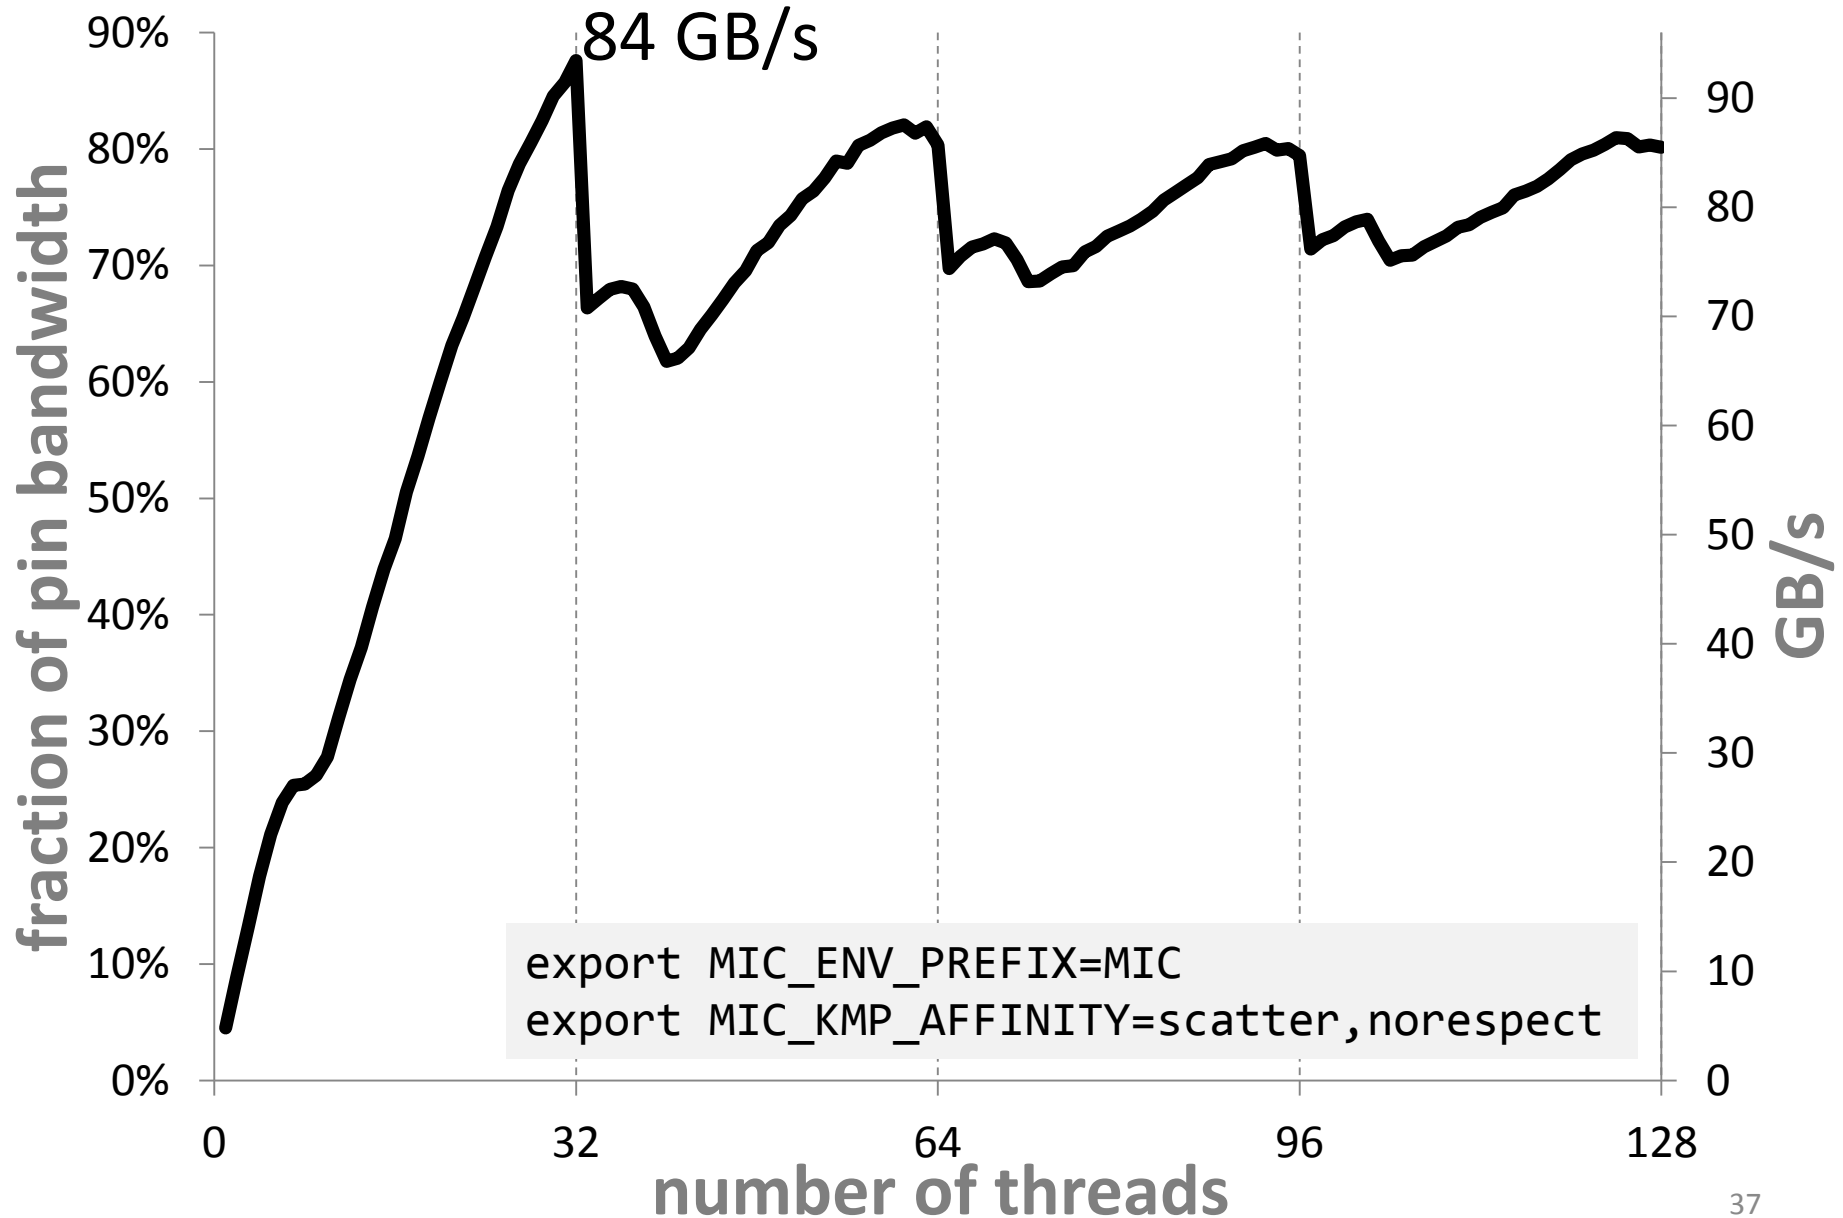
```
#include <immintrin.h>
__declspec(target(mic)) void foo( float &r, float *p, int n, int reps )
{
#ifdef __MIC__
    __m512 a = _mm512_set_1to16_ps( 0 );
    for( int j = 0; j < reps; j++ )
#pragma noprefetch
        for( int i = 0; i < n; i += 16 )
        {
            a = _mm512_add_ps( a, _mm512_load( p+i,
                _MM_FULLUPC_NONE, _MM_BROADCAST32_NONE, _MM_HINT_NONE ));
            _mm_vprefetch1( p+i+16*16, _MM_PFHINT_NONE );
            _mm_vprefetch2( p+i+32*16, _MM_PFHINT_NONE );
        }
        r = _mm512_reduce_add_ps( a );
#endif
}
```

# Manual prefetching example (2/2)

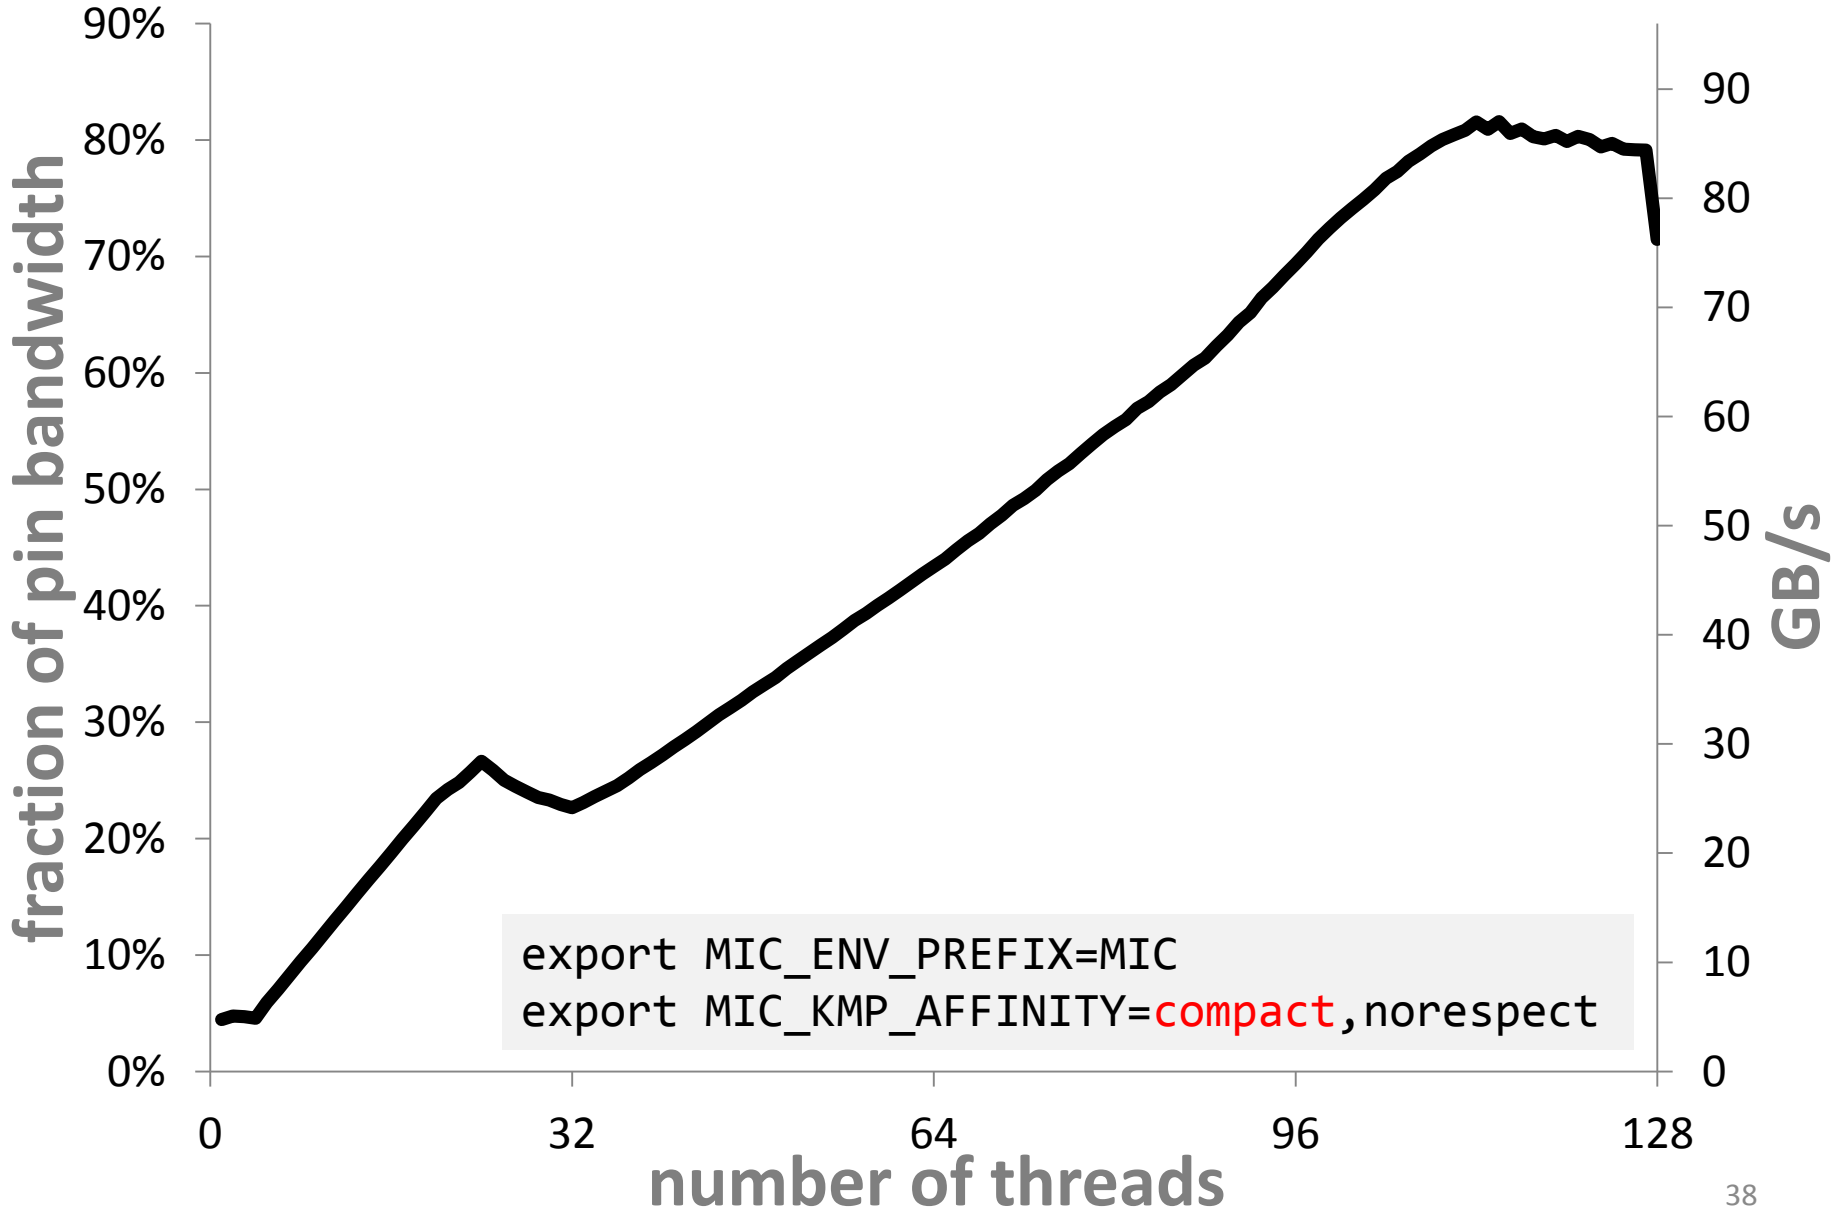
```
#include <stdio.h>
#include <omp.h>
#include <mk1.h>
main( )
{
    int n = 2<<18, nthreads_max = 128, reps = 32;
    float *A = (float*)malloc((nthreads_max*n+32*16)*sizeof(float)), r[128];
    for( int nthreads = 1; nthreads <= nthreads_max; nthreads++ )
    {
        double t;
        #pragma offload target(mic) in(A:length(nthreads*n+32*16) align(64))
        for( int k = 0; k < 3; k++ )
        {
            t = dsecnd();
            #pragma omp parallel num_threads(nthreads)
            foo( r[omp_get_thread_num()], A+n*omp_get_thread_num(), n, reps );
            t = dsecnd()-t;
        }
        double GBs = 1e-9*sizeof(float)*reps*n*nthreads / t;
        printf( "%d threads, %.3f GB/s, %g\n", nthreads, GBs, r[0] );
    }
}
```



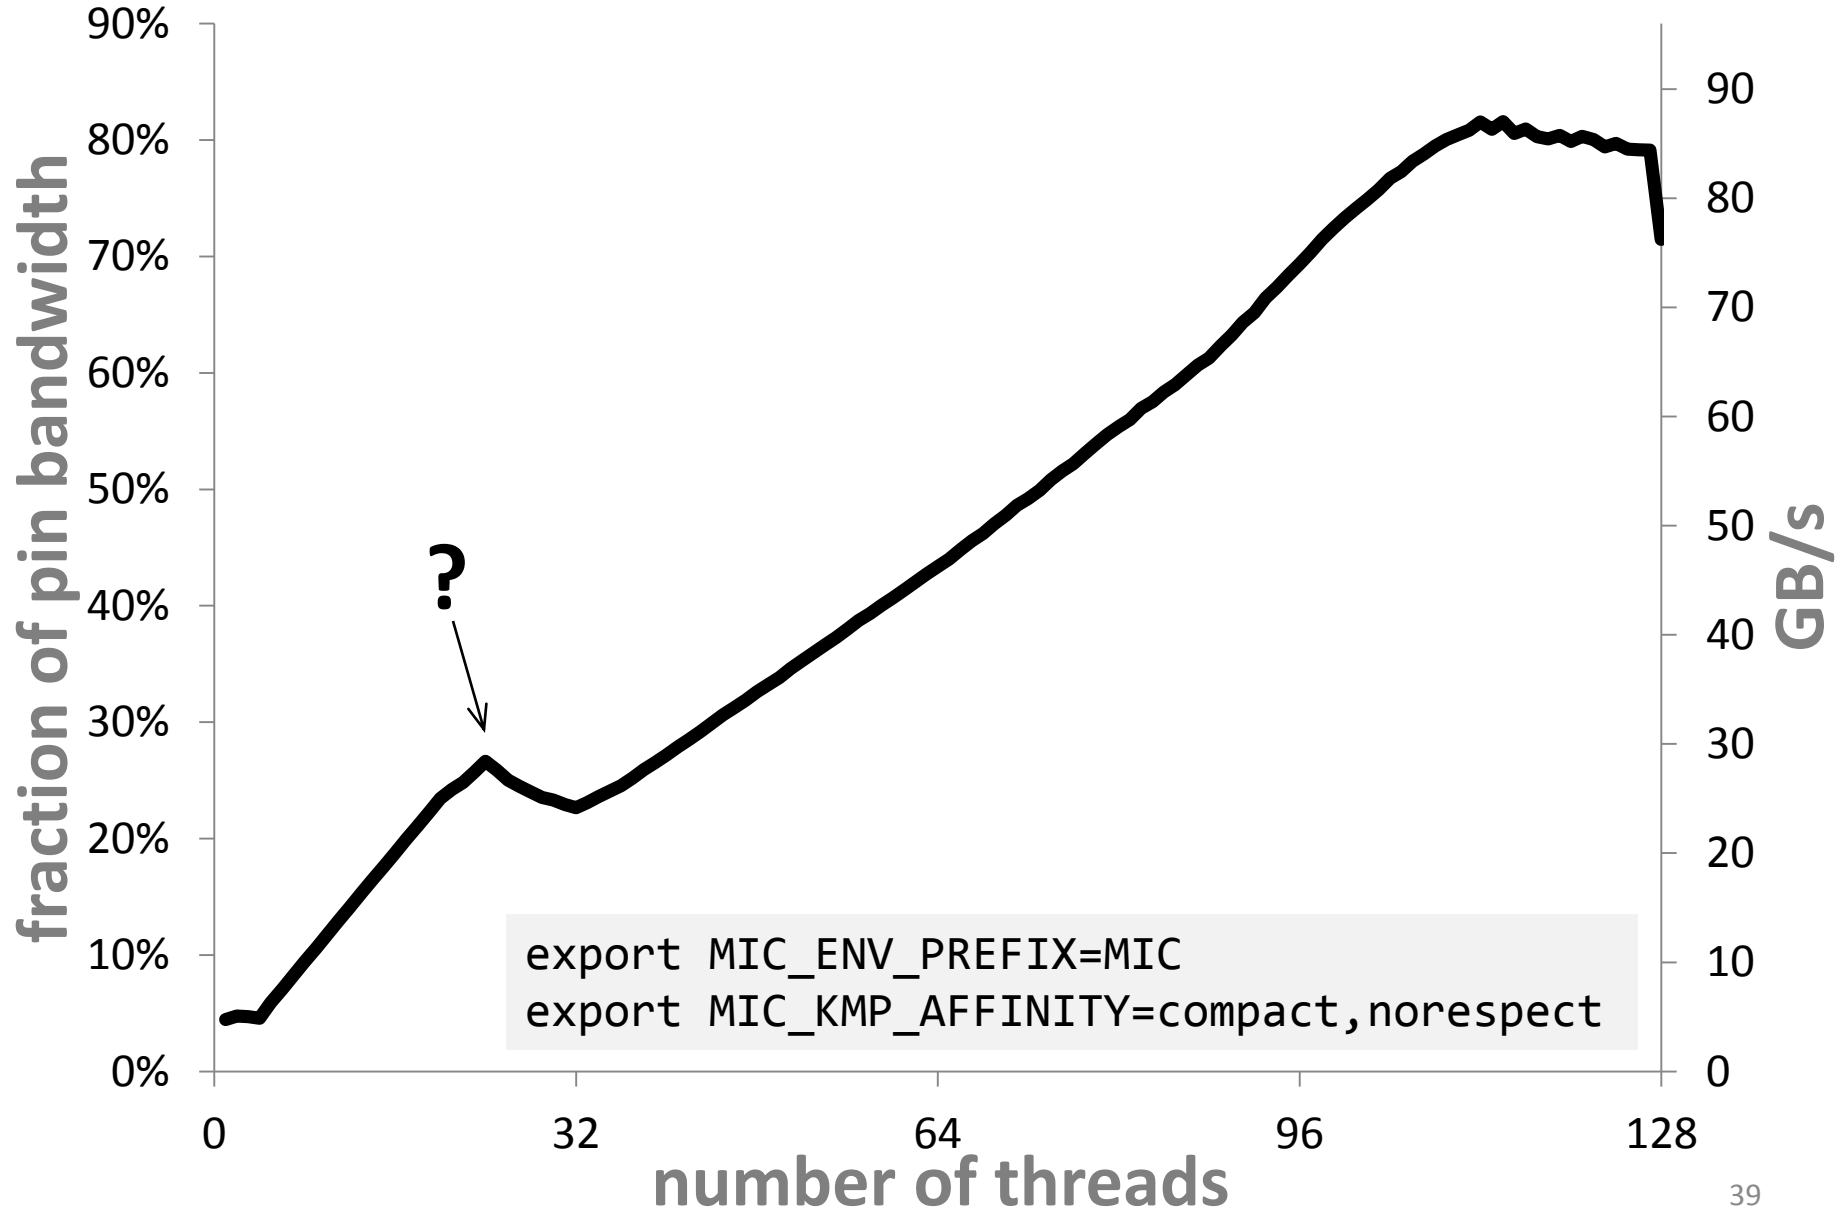
# All cores, 1 thread per core: 88% of peak



# But not if many threads, a few cores




# How would you explain this feature?




# Hypothesis: contention on ring

When assigning threads linearly, most memory and coherency traffic has to go via **2 links**:




Solution: scatter threads more evenly around the die




# Affinity settings to resolve the contention

```
export MIC_ENV_PREFIX=MIC
```

```
export MIC_KMP_AFFINITY=explicit,norespect,proclist=[1,65,33,  
97,17,81,49,113,9,73,41,105,25,89,57,121,5,69,37,101,21,85,53,  
117,13,77,45,109,29,93,61,0,2,66,34,98,18,82,50,114,10,74,42,  
106,26,90,58,122,6,70,38,102,22,86,54,118,14,78,46,110,30,94,  
62,125,3,67,35,99,19,83,51,115,11,75,43,107,27,91,59,123,7,71,  
39,103,23,87,55,119,15,79,47,111,31,95,63,126,4,68,36,100,20,  
84,52,116,12,76,44,108,28,92,60,124,8,72,40,104,24,88,56,120,  
16,80,48,112,32,96,64,127]
```



# Result: up to 1.5x speedup



# Part V

## MKL SGEMM

# Recommended settings for MKL SGEMM

MIC supports 4KB (default), 64KB and 2MB pages

Must use 2MB pages for MKL, compact affinity:

```
export MIC_ENV_PREFIX=MIC
```

```
export MIC_USE_2MB_BUFFERS=1K
```

```
export MIC_KMP_AFFINITY=compact
```

Also:


- 64B alignment
- size multiple of 16 (but 64 is slightly better)
- padding



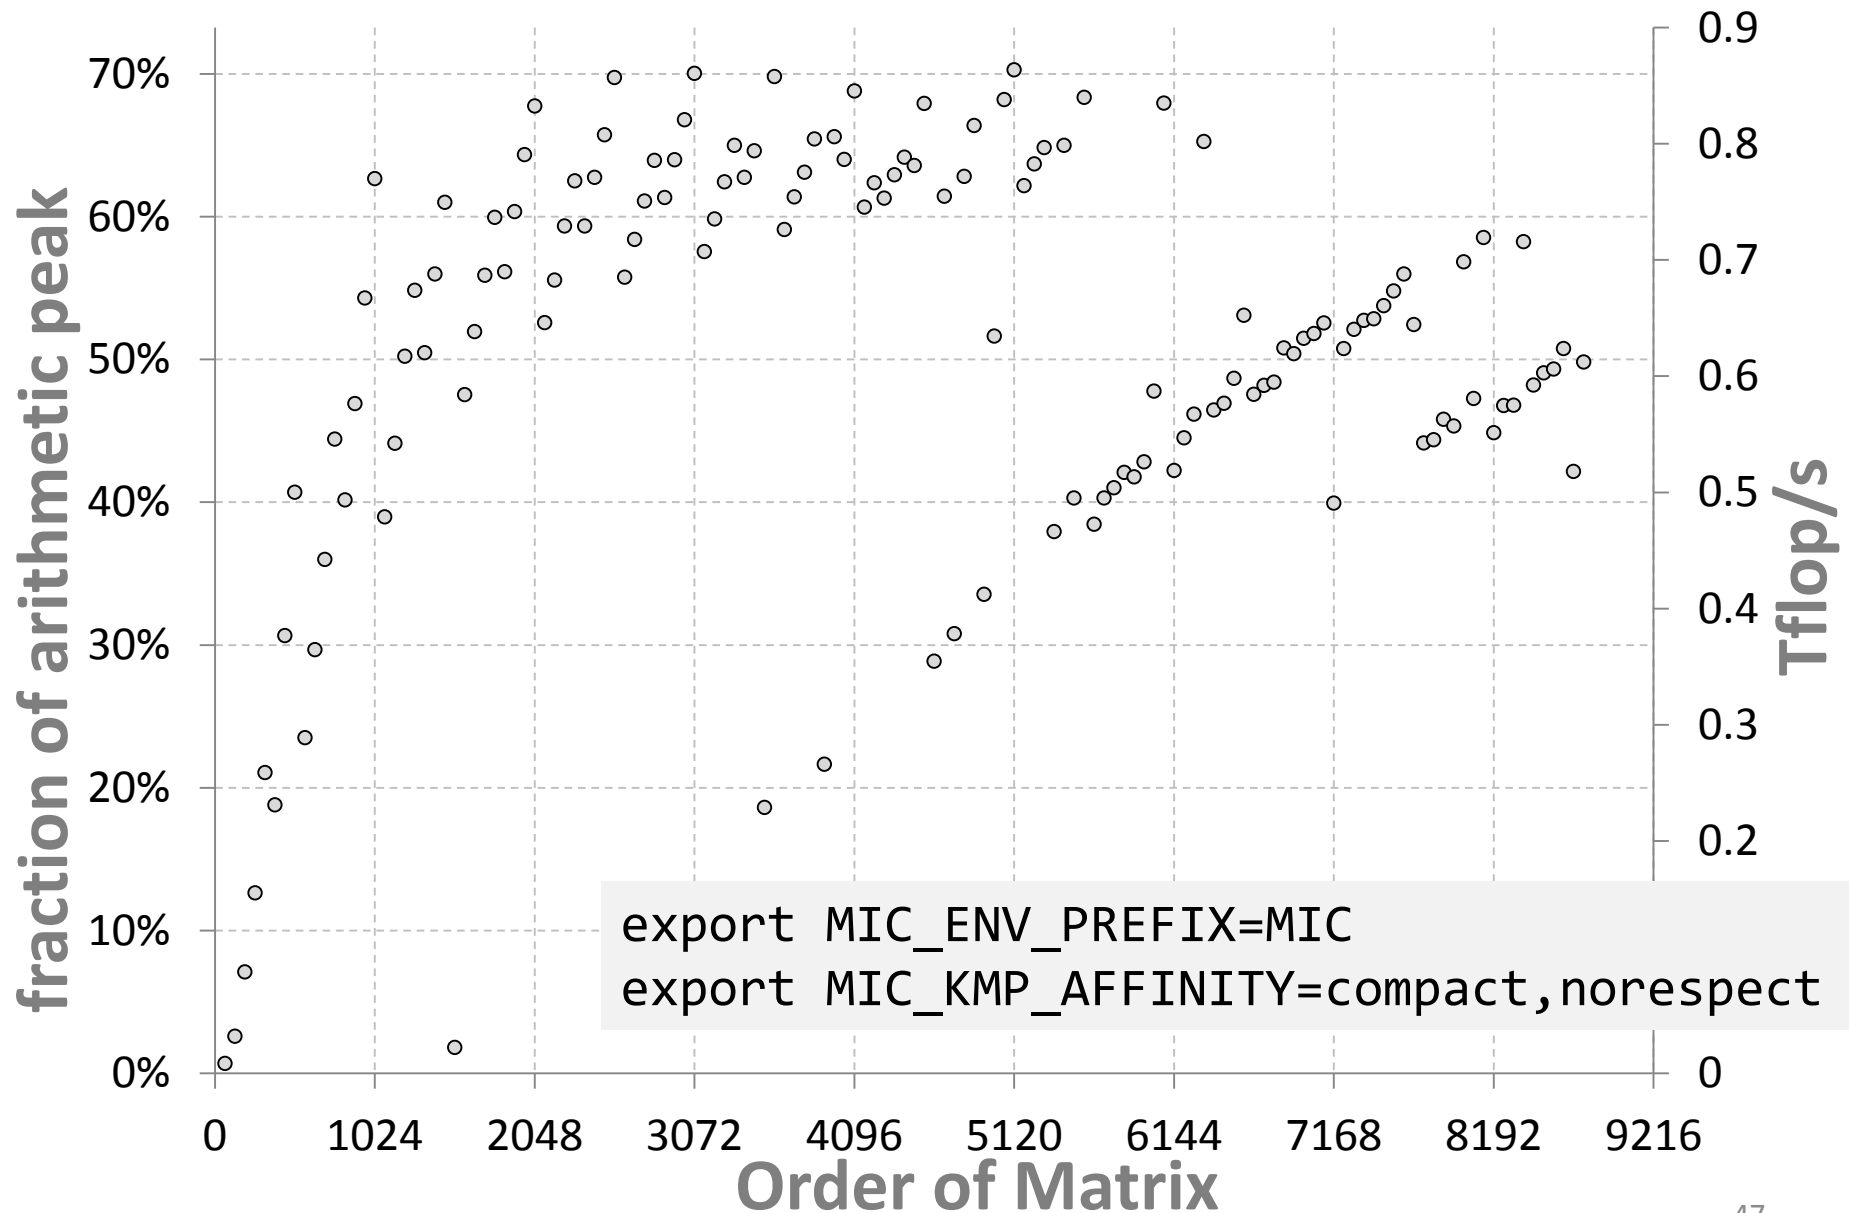
# Benchmarking SGEMM

```
#include <stdio.h>
#include <mk1.h>
#include <mk1_blas.h>
main( )
{
    int nmax = 8768;
    float *pool = (float*)malloc( 3*nmax*(nmax|64)*sizeof(float) );
    for( int n = 64; n <= nmax; n += 64 )//size multiple of 64
    {
        int lda = n|64;//a simple padding
        double t;
        for( int j = 0; j < 3*n*lda; j++ ) pool[j] = drand48( )*2-1;
        #pragma offload target(mic) inout(pool:length(3*n*lda) align(64))
        {
            float alpha=1, beta=1, *A=pool, *B=pool+n*lda, *C=pool+2*n*lda;
            sgemm( "N","N", &n,&n,&n, &alpha, A, &lda, B, &lda, &beta, C, &lda );
            t = dsecnd();
            sgemm( "N","N", &n,&n,&n, &alpha, A, &lda, B, &lda, &beta, C, &lda );
            t = dsecnd() - t;
        }
        printf( "%d %g Gflop/s \n", n, 2e-9*n*n*n/t );
    }
}
```

# Performance of SGEMM on 31 cores



# Performance of SGEMM on 32 cores



# Benchmarking SGEMM in native mode

```
#include <stdio.h>
#include <mk1.h>
#include <mk1_blas.h>
#include <sys/mman.h> // for allocating 2 MB pages
main( )
{
    int nmax = 8768;
    float *pool = (float*)mmap( NULL, 3*nmax*(nmax|64)*sizeof(float),
PROT_READ | PROT_WRITE, MAP_ANONYMOUS | MAP_PRIVATE | MAP_HUGETLB, -1, 0 );
    for( int n = 64; n <= nmax; n += 64 )
    {
        int lda = n|64;
        double t;
        for( int j = 0; j < 3*n*lda; j++ ) pool[j] = drand48( )*2-1;
        float alpha=1, beta=1, *A=pool, *B=pool+n*lda, *C=pool+2*n*lda;
        sgemm( "N","N", &n,&n,&n, &alpha, A, &lda, B, &lda, &beta, C, &lda );
        t = dsecnd();
        sgemm( "N","N", &n,&n,&n, &alpha, A, &lda, B, &lda, &beta, C, &lda );
        t = dsecnd() - t;
        printf( "%d %g Gflop/s\n", n, 2e-9*n*n*n/t );
    }
}
```

# Running benchmark in native mode

Host: compile, upload the binary and libiomp5.so

```
icc -mmic -O3 -mk1 bench.cpp
```

```
sudo micput 192.168.1.100 ./a.out /tmp/a.out
```

```
sudo micput 192.168.1.100 /opt/intel/composerxe_mic/  
    compiler/lib/mic/libiomp5.so /tmp
```

```
sudo ssh -i /opt/intel/mic/id_rsa root@192.168.1.100
```

MIC: configure 2MB page allocation and affinity

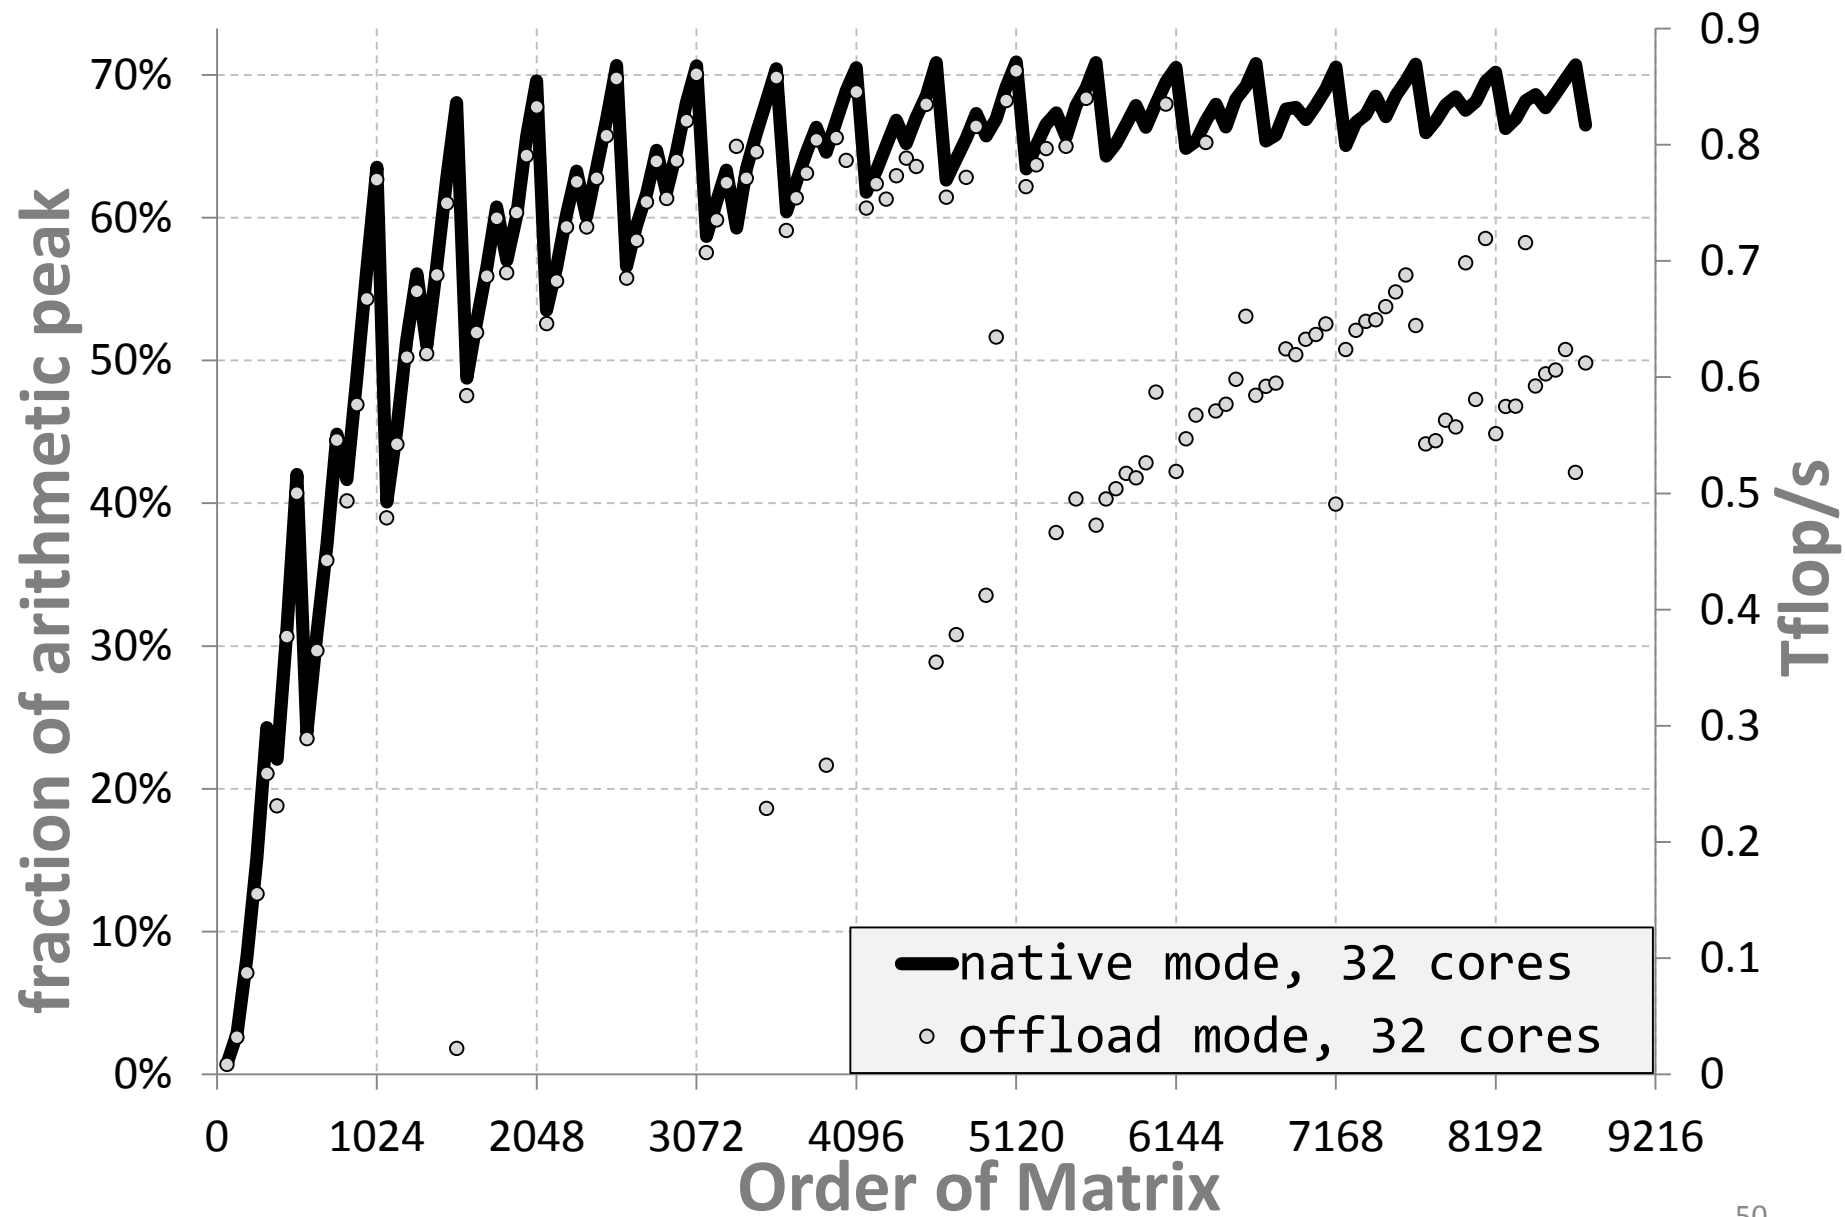
```
echo 500 > /proc/sys/vm/nr_hugepages
```

```
export KMP_AFFINITY=compact
```

```
export LD_LIBRARY_PATH=/tmp
```

```
/tmp/a.out
```

# Using 32 cores: native vs offload mode



## Part VI

MKL SGETRF ( $\approx$ LINPACK benchmark)

# Recommended settings for SGETRF

Similar to SGEMM:

- 64B alignment
- Padding
- Size multiple of 64

Allocate extra trailing space (for prefetching?)

```
export MIC_ENV_PREFIX=MIC
```

```
export MIC_USE_2MB_BUFFERS=1K
```

```
export MIC_KMP_AFFINITY=compact
```




# Benchmarking SGETRF


```
#include <stdio.h>
#include <mk1.h>
#include <mk1_lapack.h>
main( )
{
    const int nmax = 14336;
    float *A = (float*)malloc( (nmax+128)*(nmax|64)*sizeof(float) );
    for( int n = 64; n <= nmax; n += 64 )//size multiple of 64
    {
        double t;
        int ipiv[nmax], lda = n|64;//padding
        #pragma offload target(mic) in(A:length((n+128)*lda) align(64))
        {
            for( int info, i = 0; i < 3; i++ )
            {
                for( int j = 0; j < n*lda; j++ ) A[j] = drand48( )*2-1;
                t = dsecnd( );
                sgetrf( &n,&n, A, &lda, ipiv, &info );
                t = dsecnd() - t;
            }
        }
        printf( "%d %g Gflop/s\n", n, 2e-9*n*n*n/3/t );
    }
}
```

extra space for prefetching?  
↙

# 52% of peak only (was 68% in SGEMM)



# Using all 32 cores is troublesome again




## **Part VII**

### **Global barrier**

# A simplistic barrier benchmark

```
#include <stdio.h>
#include <immintrin.h>
#include <omp.h>
main()
{
    for( int nthreads = 2; nthreads <= 124; nthreads++ )
    {
        double cycles;
        #pragma offload target(mic)
        {
            int n = 100000;
            __int64 c0 = _rdtsc();
            #pragma omp parallel num_threads(nthreads)
            {
                for( int i = 0; i < n; i++ )
                {
                    #pragma omp barrier
                }
            }
            cycles = (_rdtsc() - c0)/(double)n;
        }
        printf( "%d threads, %g cycles\n", nthreads, cycles );
    }
}
```

# 10,000 cycles per barrier?



# Barrier cost is important

N.B. 7,000 cycles  $\approx$  7,000,000 flops at peak rate

**Expensive synchronization is a common motivation for heterogeneous algorithms**

E.g. synchronization cost with GPU is  $\approx 5\mu\text{s}$ , so we typically leave synchronization-heavy work on CPU

- More expensive synchronization on a device = more work should be run elsewhere

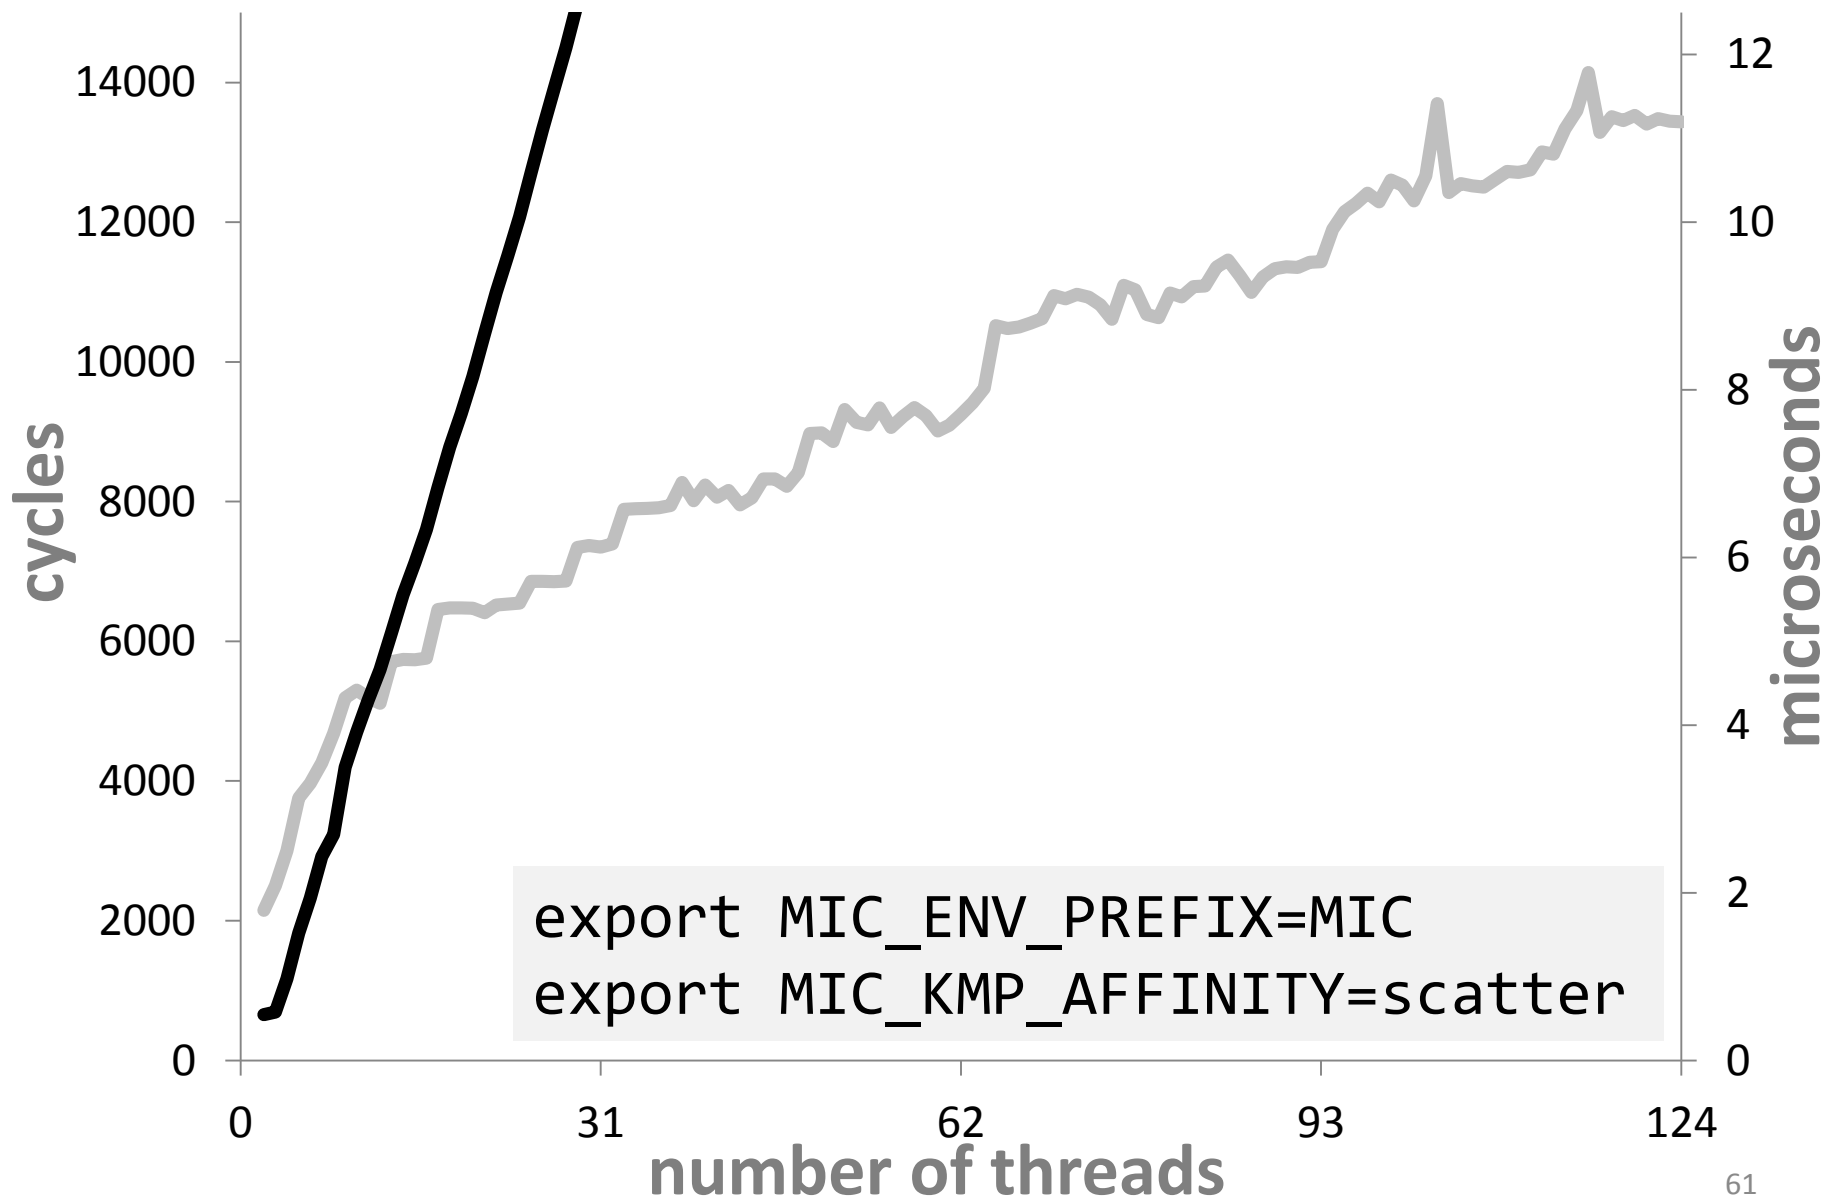
How sync across die can be as expensive as sync across PCIe?

# Naïve barrier

```
#include <stdio.h>
#include <immintrin.h>
main()
{
    for( int nthreads = 2; nthreads <= 124; nthreads++ )
    {
        double cycles;
        #pragma offload target(mic)
        {
            int n = 100000;
            volatile int counter = 0;
            __int64 c0 = _rdtsc();
            #pragma omp parallel num_threads(nthreads)
            {
                for( int i = 0; i < n; i++ )
                {
                    __sync_fetch_and_add( &counter, 1 );//vote
                    while( counter < (i+1)*nthreads ); //wait until all voted
                }
            }
            cycles = (_rdtsc() - c0)/(double)n;
        }
        printf( "%d threads, %g cycles\n", nthreads, cycles );
    }
}
```




# Naïve barrier is slower



# Simple barrier with sense

```
#include <stdio.h>
#include <immintrin.h>
main()
{
    for( int nthreads = 2; nthreads <= 124; nthreads++ )
    {
        double cycles;
        #pragma offload target(mic)
        {
            int n = 100000;
            __declspec(align(64)) volatile int counter = 0, sense = 0;
            __int64 c0 = _rdtsc();
            #pragma omp parallel num_threads(nthreads)
            {
                for( int i = 1; i <= n; i++ )
                {
                    int old = __sync_fetch_and_add( &counter, 1 );
                    if( old+1 == i*nthreads ) sense = i;
                    else while( sense < i );
                }
            }
            cycles = (_rdtsc() - c0)/(double)n;
        }
        printf( "%d threads, %g cycles\n", nthreads, cycles );
    }
}
```

# Simple barrier beats OpenMP



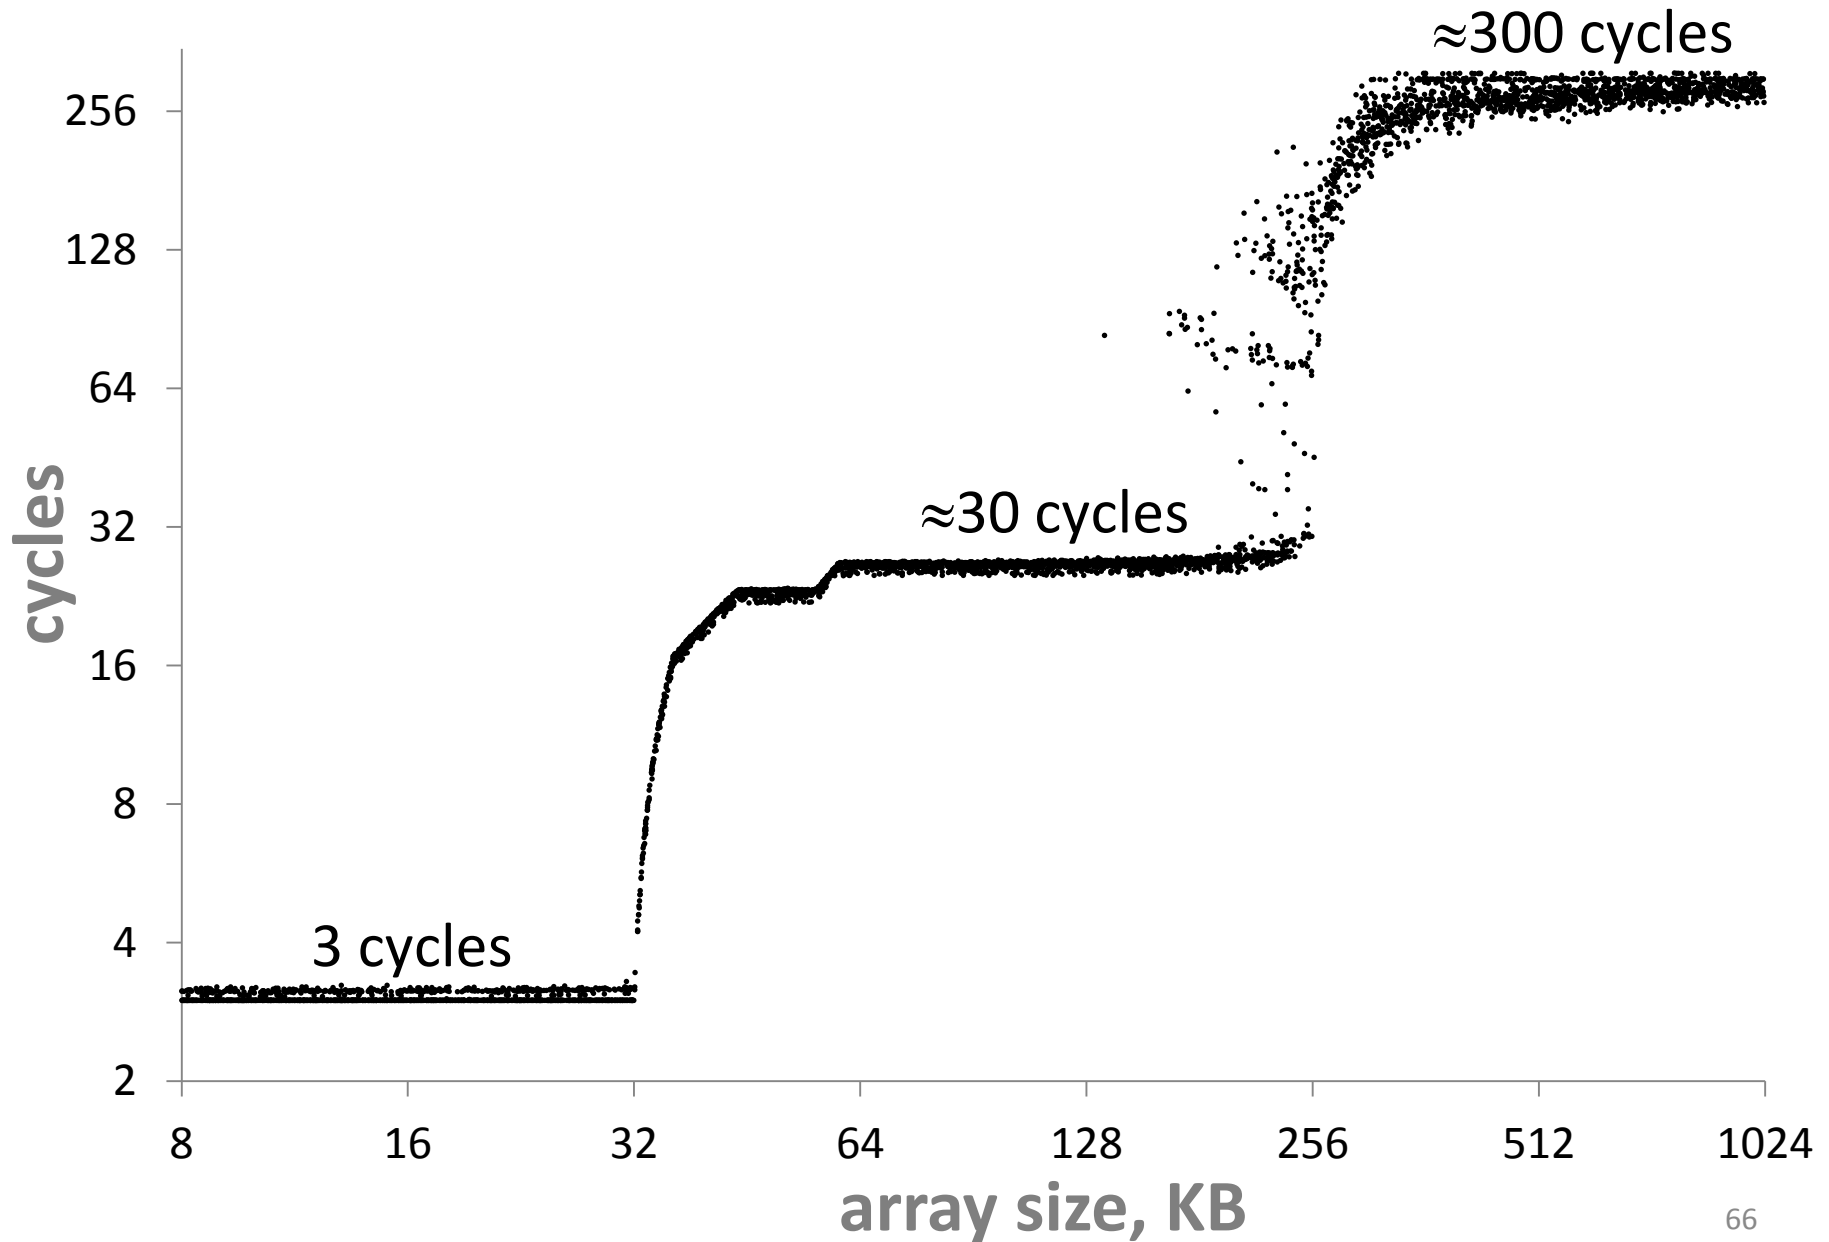
## **Part VIII**

### Memory and cache latency

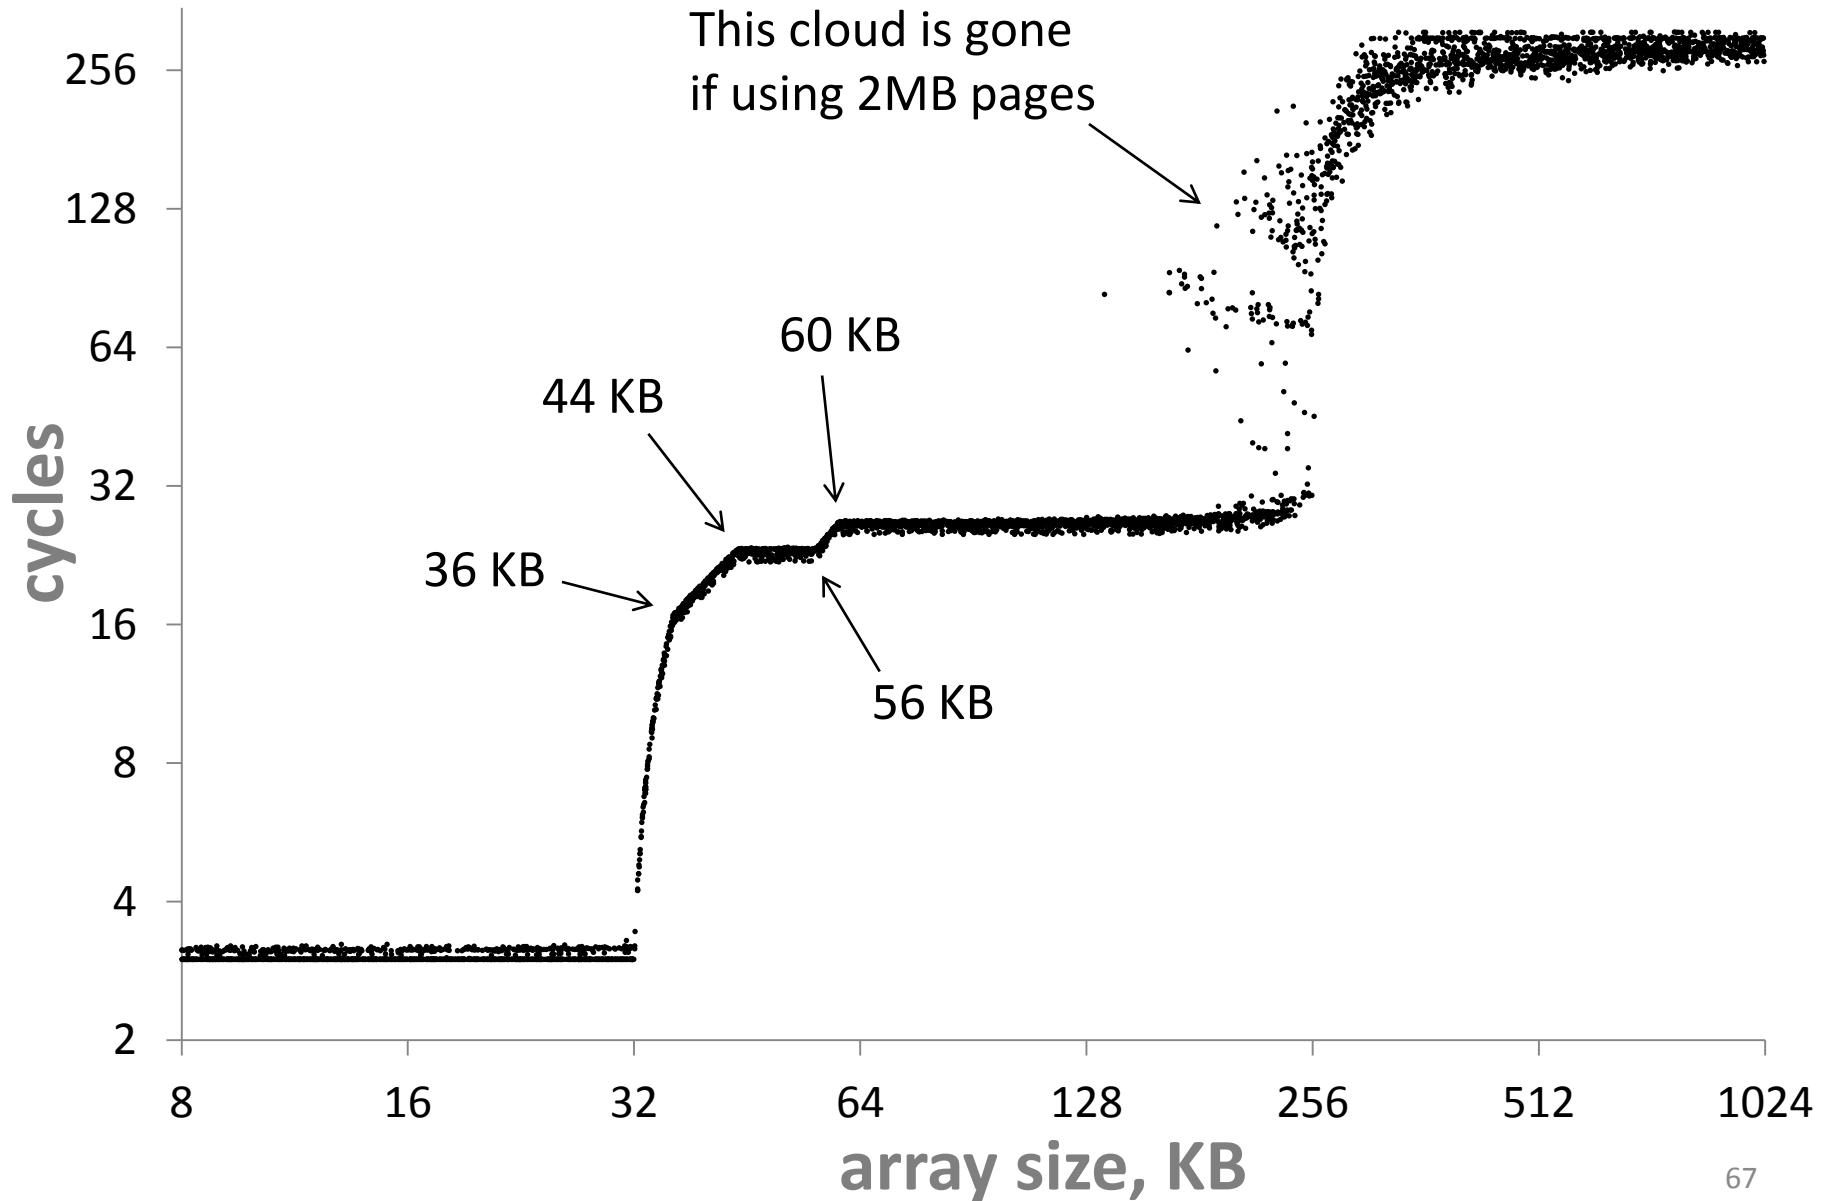
# Pointer chasing benchmark

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>
#include <immintrin.h> //this time we use time stamp counter
main( )
{
    int nmin = 1<<10, nmax = 1<<17, reps = 1<<20;
    long *next = (long*)malloc( nmax*sizeof(long) );
    for( int i = 0; i < 5000; i++ )
    {
        int n = exp( log(nmin) + drand48()*(log(nmax)-log(nmin)));
        for( int i = 0; i < n; i++ ) next[i] = (i+8)%n; //stride = cache line
        float c;
        long j = lrand48()%n;
        #pragma offload target(mic) in(next:length(n))
        {
            for( int it = 0; it < n; it++ ) j = next[j]; //warm up
            __int64 t = _rdtsc();
            for( int it = 0; it < reps; it += 2 ) j = next[next[j]]; //run
            c = (float)( _rdtsc() - t ) / reps;
        }
        printf( "%.3f KB, %.3f cycles, %ld\n", n*sizeof(long)/1024.f, c, j );
    }
}
```

# Classical geometric latency distribution



# Fancy L2 structure @ multiples of 4KB

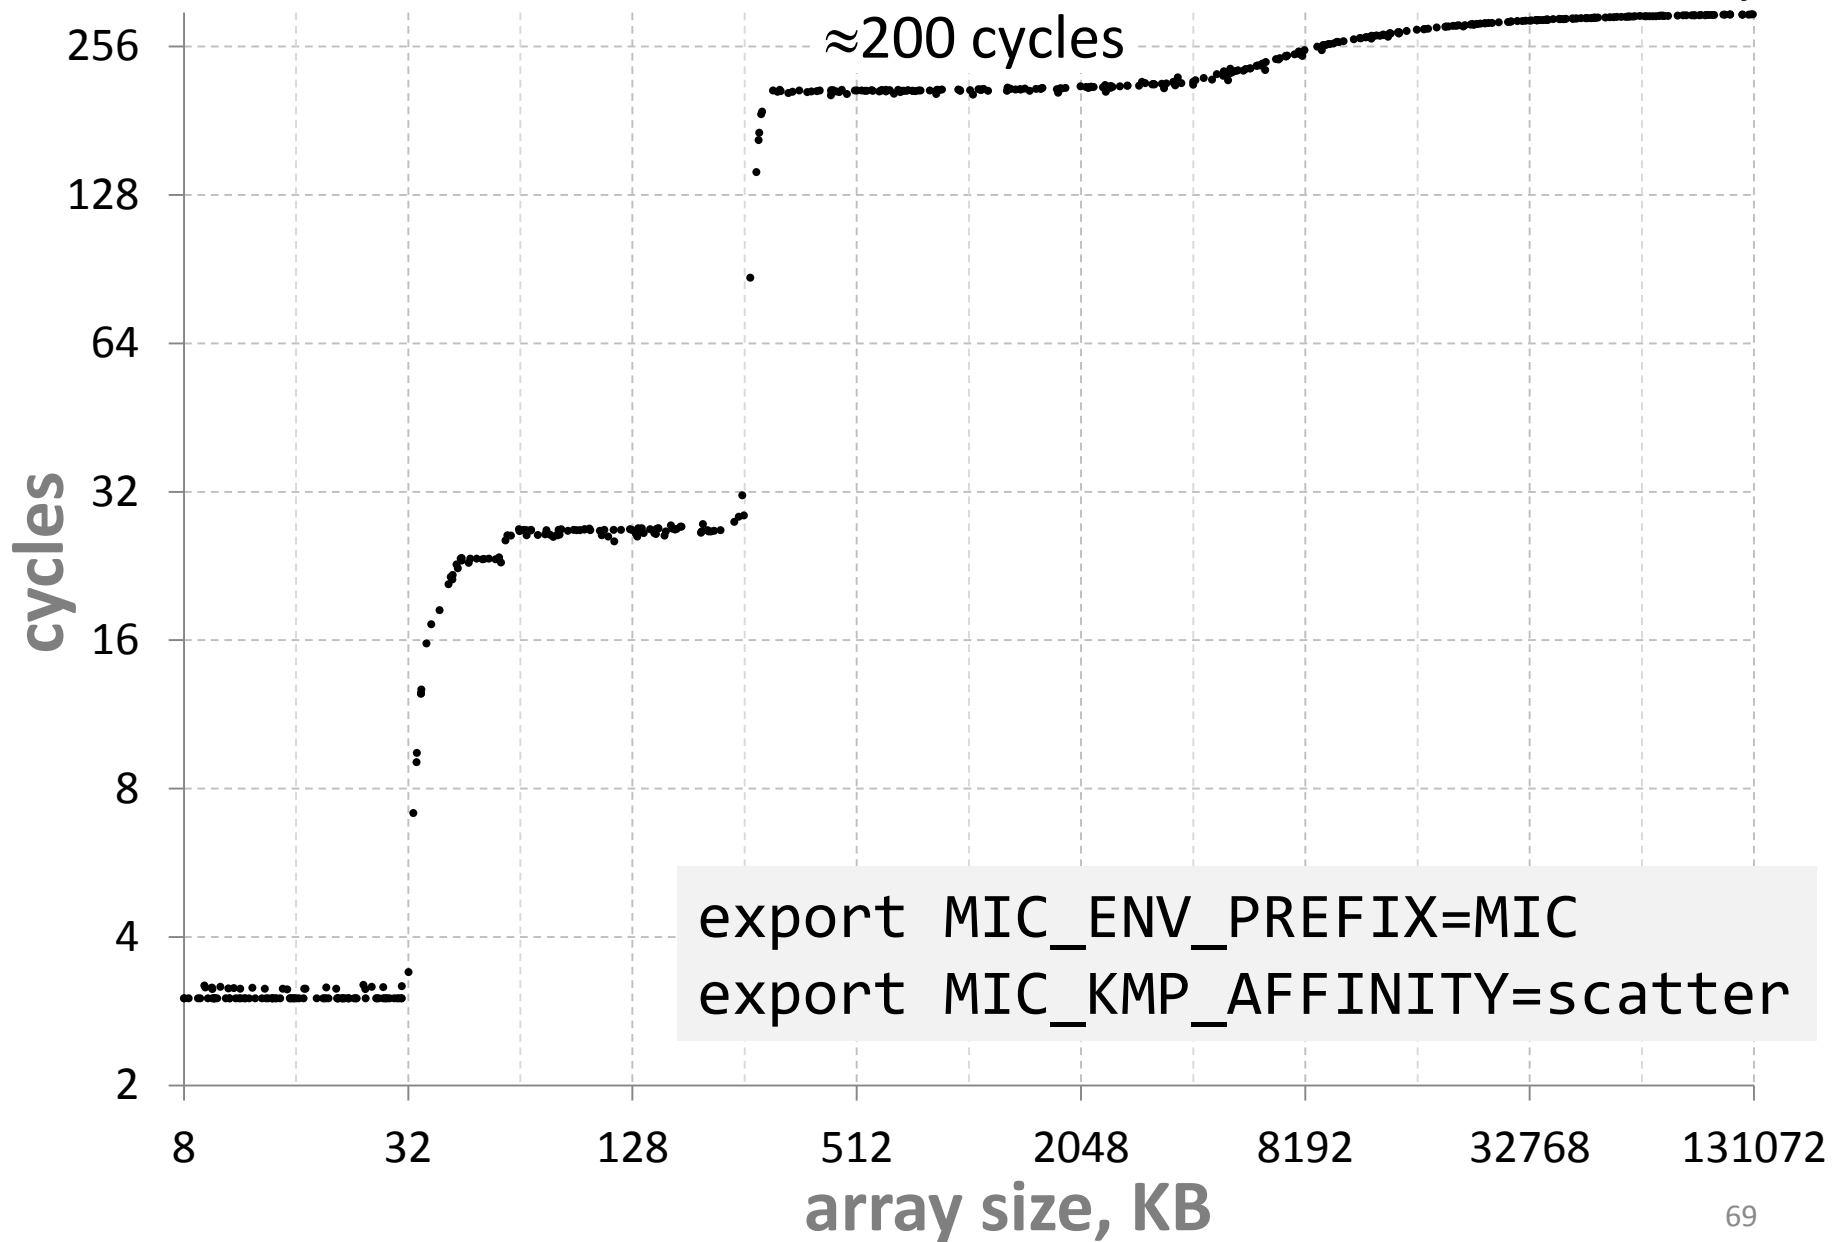


# Put array into distributed L2

```
#include <omp.h>
#include <math.h>
#include <stdio.h>
#include <immintrin.h>
main( ) {
    int nmin = 1<<10, nmax = 1<<24, reps = 1<<20, nthreads = 31;
    long *next = (long*)malloc( nmax*sizeof(long) );
    for( int i = 0; i < 500; i++ )
    {
        float cycles;
        int n = exp( log(nmin) + drand48()*(log(nmax)-log(nmin)));
        long j = lrand48()%n;
        #pragma offload target(mic) in(next:length(n))
            #pragma omp parallel num_threads(nthreads)
            {
                int tid = omp_get_thread_num(), block = (n+nthreads-2)/(nthreads-1);
                if( tid > 0 )
                    for( int i = (tid-1)*block; i < tid*block && i < n; i++ )
                        next[i] = (i+8)%n; //hope to get it in distributed L2
                #pragma omp barrier
                if( tid == 0 ) //thread 0 runs the bench
                {
                    for( int it = 0; it < n; it++ ) j = next[j]; //warm up
                    __int64 t = _rdtsc();
                    for( int it = 0; it < reps; it += 2 ) j = next[next[j]]; //bench
                    cycles = (float)( _rdtsc() - t ) / reps;
                }
            }
        printf( "%.3f KB %.3f cycles %ld\n", n*sizeof(long)/1024.f, cycles, j );
    }
}
```



# Remote L2 is about as far as DRAM



## **Part IX**

### Ping-pong latency


# Ping-pong benchmark (1/2)

```
#include <stdio.h>
#include <omp.h>
#include <immintrin.h>

__declspec(target(mic)) void ping ( volatile int *p, int n )
{
    for( int i = 0; i < n; i += 2 )
    {
        *p = i+1;           // send a signal to slave
        while( *p != i+2 ); // wait for reply
    }
}

__declspec(target(mic)) void pong( volatile int *p, int n )
{
    for( int i = 0; i < n; i += 2 )
    {
        while( *p != i+1 ); // wait for a signal from master
        *p = i+2;           // reply
    }
}
```

# Coherency triangle



Directory is distributed over 32 stations (DTD)

Station number is found as hashed cache line #


Q: how latency depends on the choice of station?

```

main() {
    const int nslots = 16*32, reps = 1000, ncores = 32;
    for( int th1, th2, j = 0; j < ncores*ncores; j++ )
    {
        if( (th1=j%ncores) == (th2=j/ncores) ) continue;
        __int64 t, tmin, tavg = 0, tmax;
        #pragma offload target(mic)
        {
            volatile int slots[nslots][16]; //try many different DTDs
            for( int slot = 0; slot < nslots; slot++ )
            {
                slots[slot][0] = 0;
                #pragma omp parallel num_threads(ncores)
                for( int k = 0; k < 3; k++ )
                {
                    if( omp_get_thread_num() == th1 )
                    {
                        __int64 t0 = _rdtsc();
                        ping( &slots[slot][0], reps );
                        t = t < (t0=_rdtsc()-t0) && k ? t : t0;
                    }
                    else if( omp_get_thread_num() == th2 )
                    {
                        pong( &slots[slot][0], reps );
                    }
                }
                tavg += t;
                tmin = slot == 0 || t < tmin ? t : tmin;
                tmax = slot == 0 || t > tmax ? t : tmax;
            }
        }
        printf( "%d %d %g %g %g\n", th1, th2, tmin/(float)reps,
                tavg/(float)reps/nslots, tmax/(float)reps );
    }
}

```

# Worst vs Best: 1.4x gap



# Average across all samples: 370 cycles

Compare to 200 cycle latency to remote L2 slice

Why different?

- Each ping requires 2 directory accesses
- First for reading, second for writing:

```
for( int i = 0; i < n; i += 2 )
{
    while( *p < i+1 ); // get shared rights
    *p = i+2;          // get exclusive rights
}
```

Can do using only 1 directory access!

- Get exclusive rights when spinning


# New ping-pong kernels

```
#define hint _MM_PFHINT_EX
__declspec(target(mic)) void ping( volatile int *p, int n )
{
#ifdef __MIC__
    for( int i = 0; i < n; i += 2 )
    {
        *p = i+1;
        do { _mm_vprefetch1((const void*)p, hint); } while( *p != i+2 );
    }
#endif
}

__declspec(target(mic)) void pong( volatile int *p, int n )
{
#ifdef __MIC__
    for( int i = 0; i < n; i += 2 )
    {
        do { _mm_vprefetch1((const void*)p, hint); } while( *p != i+1 );
        *p = i+2;
    }
#endif
}
```



# Result: uniform 2.0x-2.3x speedup




# Ping-pong timing summary

	Best core choice	Worst core choice
Best DTD choice	110	160
Average	160	180
Worst DTD choice	190	220

**Up to 2x speedup by optimizing both core and directory placement**

# Best ping time, core A is #9 (corner)



## **Part X**

### Broadcasting shared variable

# Broadcast benchmark (1/2)

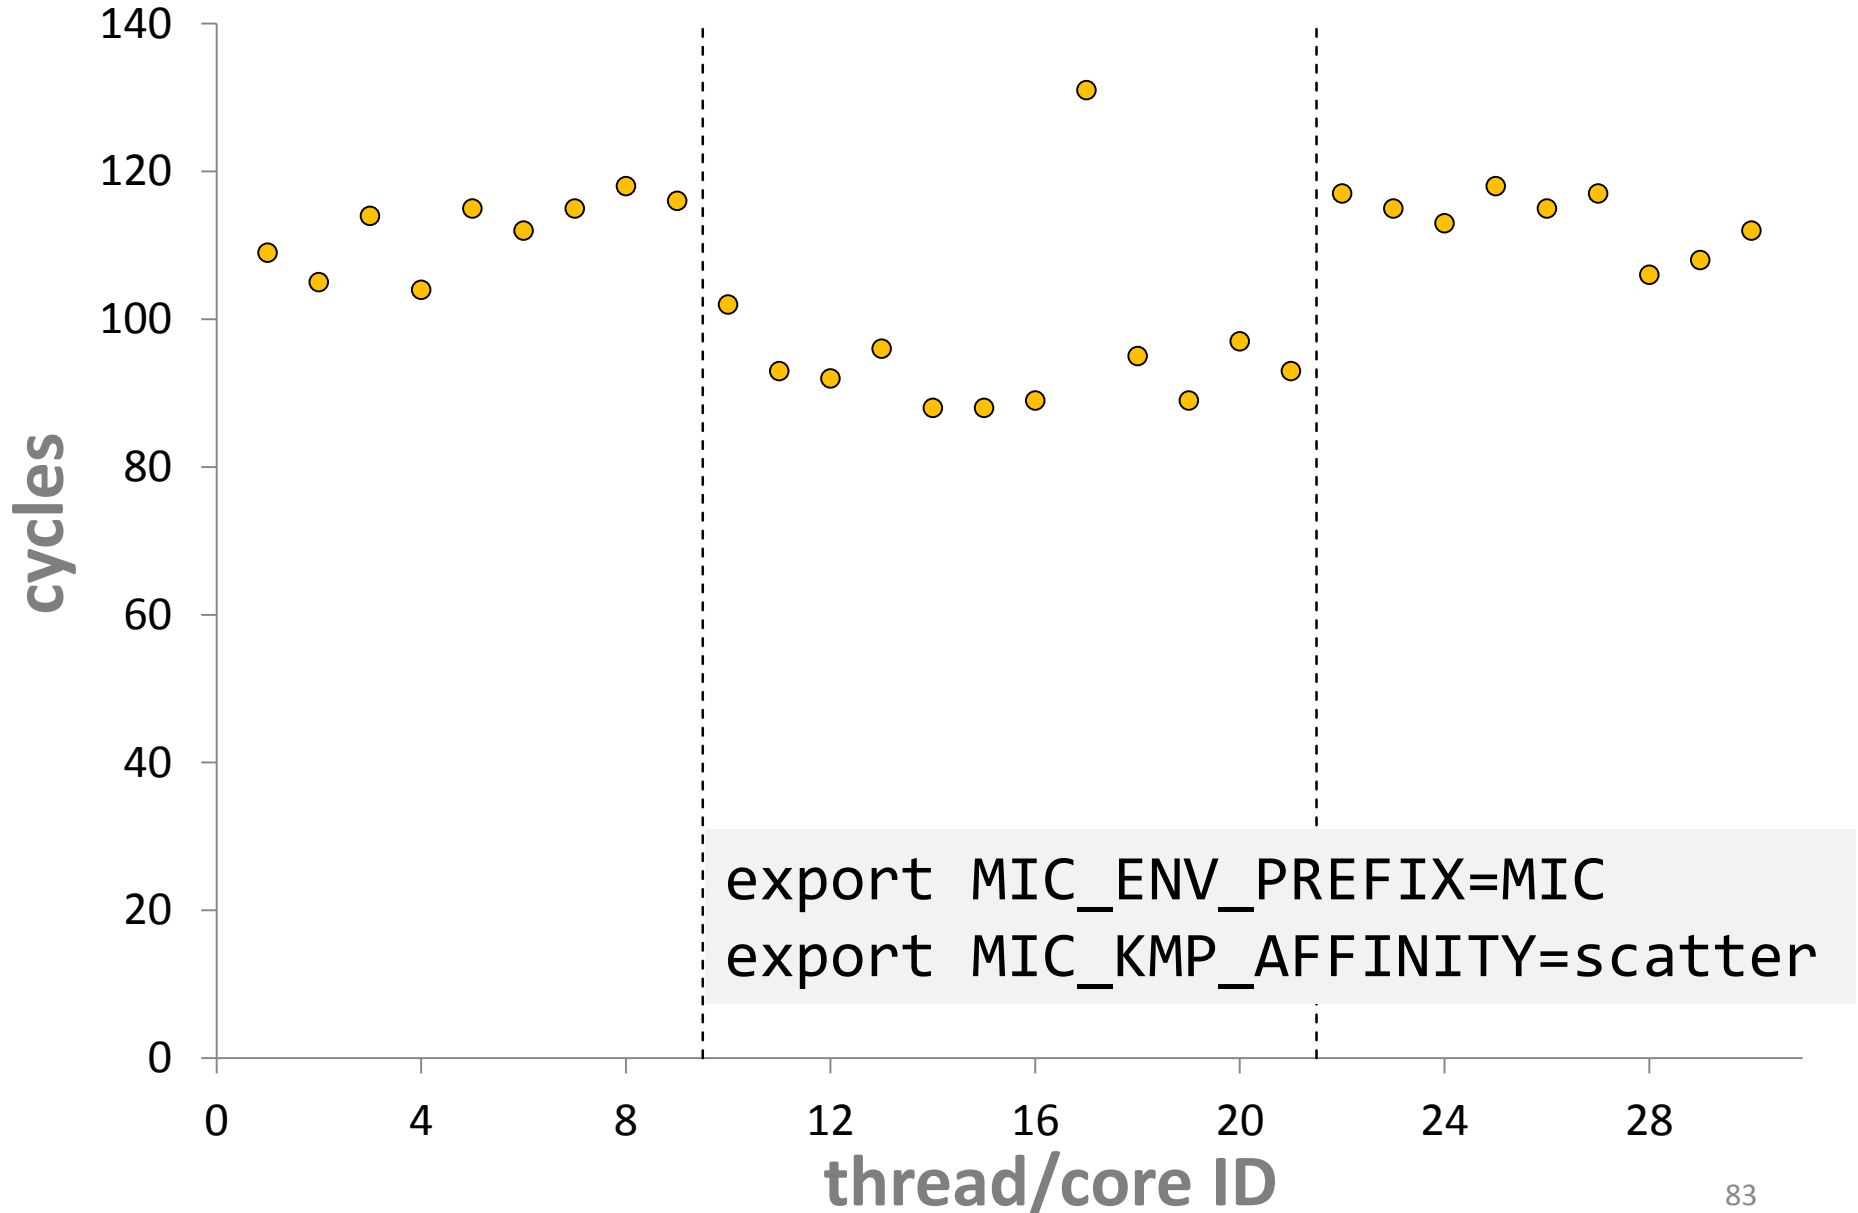
```
#include <mkl.h>
#include <immintrin.h>
__declspec(target(mic)) void master( __int64 *r, volatile int &s, int n )
{
    for( int i = 0; i < n; i++, r += 2 )
    {
        for( double t = dsecnd(); dsecnd() - t < 1e-3; );//pause
        r[0] = _rdtsc(); //time before
        s = i;           //send signal
        r[1] = _rdtsc(); //time after
    }
}

__declspec(target(mic)) void slave( __int64 *r, volatile int &s, int n )
{
    for( int i = 0; i < n; i++, r += 2 )
    {
        do { r[0] = _rdtsc(); } //time before
        while( s != i );        //receive signal
        r[1] = _rdtsc();        //time after
    }
}
```


# Broadcast benchmark (2/2)

```
#include <omp.h>
#include <stdio.h>
main()
{
    const int nthreads = 31, n = 4;
    __int64 r[nthreads][n][2];
    #pragma offload target(mic)
    {
        volatile int sense = -1;
        #pragma omp parallel num_threads(nthreads)
        {
            int tid = omp_get_thread_num();
            if( tid == 0 ) master( &r[tid][0][0], sense, n ); //one master
            else                 slave( &r[tid][0][0], sense, n ); //many slaves
        }
    }
    for( int t = 0; t < nthreads; t++ )
    {
        for( int i = 0; i < n; i++ ) printf( "%5ld", r[t][i][0]-r[0][i][0] );
        for( int i = 0; i < n; i++ ) printf( "%5ld", r[t][i][1]-r[t][i][0] );
        printf( "\n" );
    }
}
```

# Time until invalidation request is received

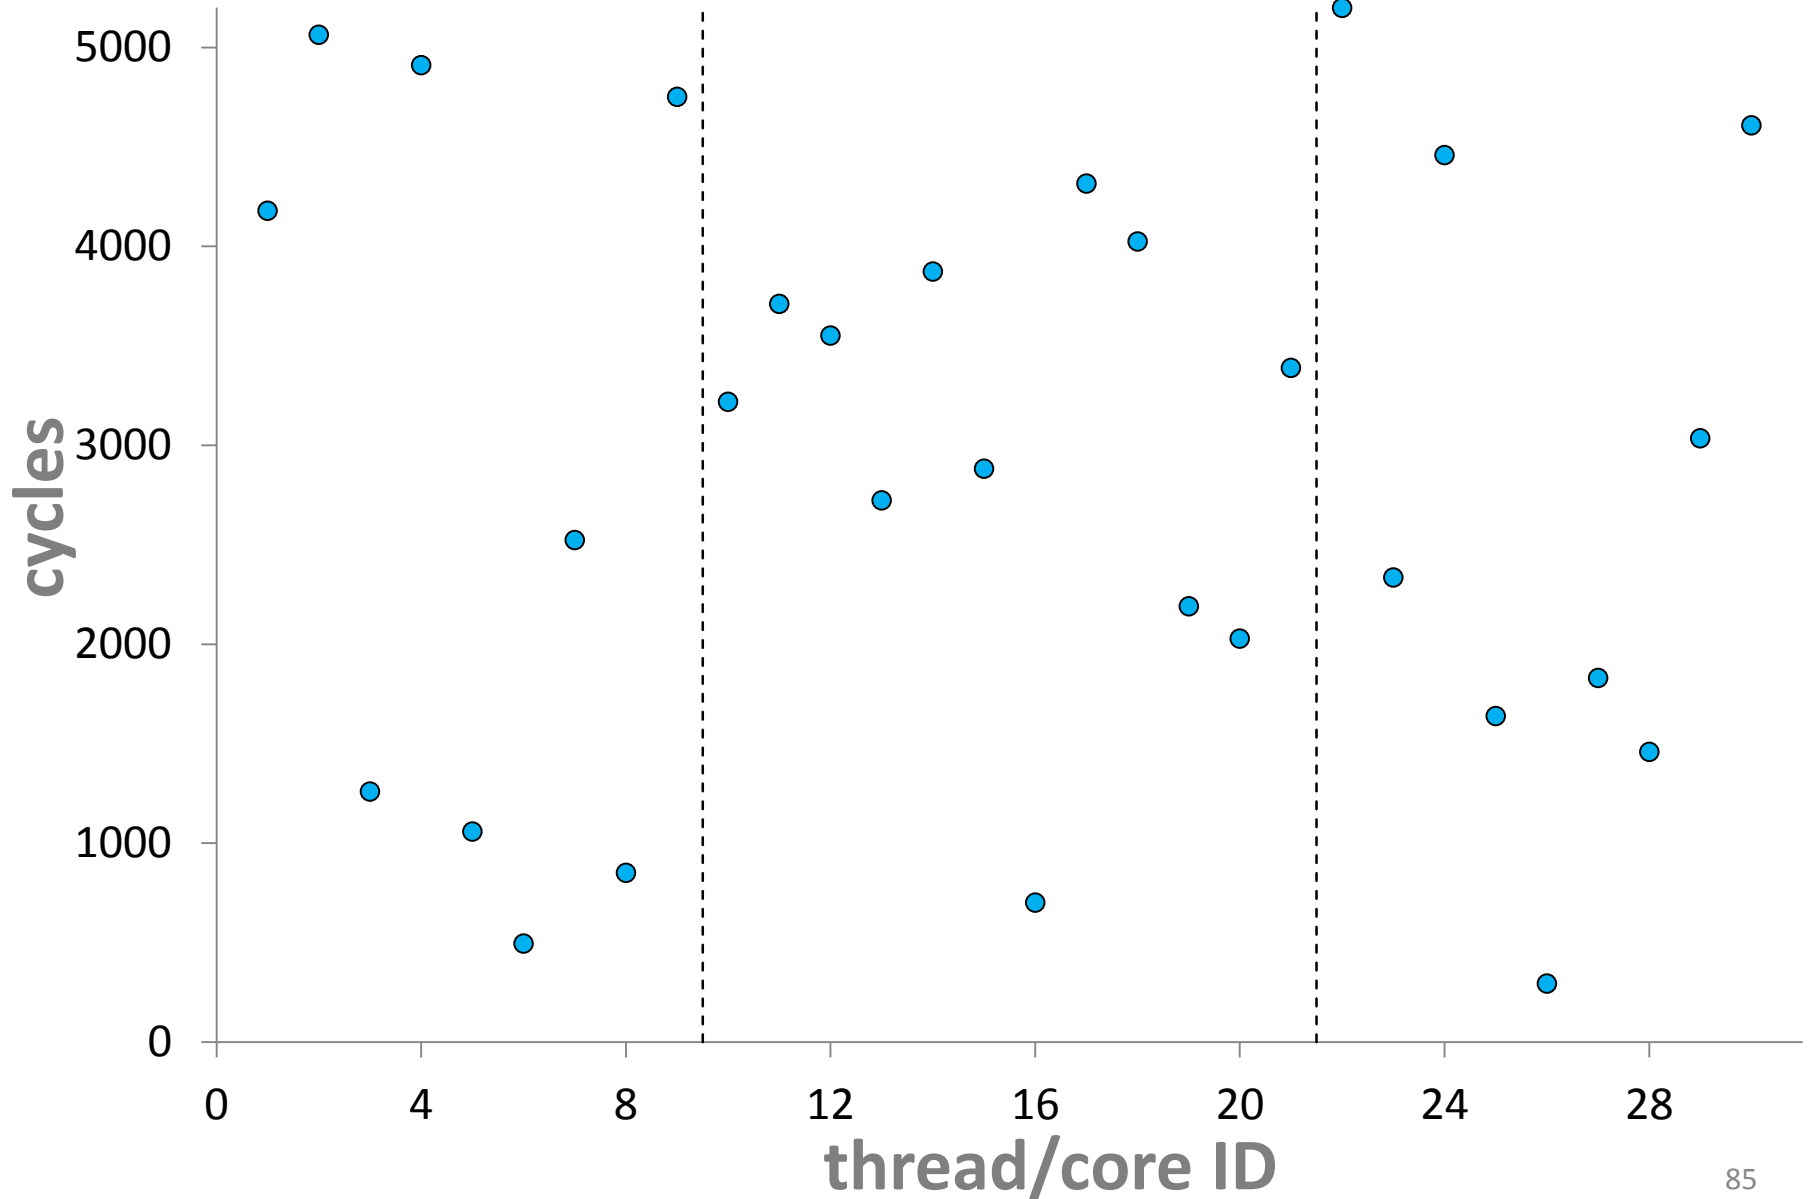


# Time until invalidation, sorted







# Delay until updated copy is retrieved



# Delay until the update, sorted



# Master and slave timings together



# Programming MIC is fun!

Thanks to Kostadin Ilov for technical support; Roman Dubtsov for help with MKL; Kevin Davis for help on offload overhead; James Cownie for help on OpenMP barrier; Timothy Prince for affinity settings; Charles Congdon for shared virtual memory explanation; Romain Dolbeau for Intel memory consistency model